

- <https://github.com/timestocome> (<https://github.com/timestocome>)

This problem is just a slight variation on MNIST. It's a different number set and there is a labeled training set, a labeled holdout set and an unlabeled set for submitting to Kaggle.

What makes it interesting is that all 3 sets of integers have different std, means

The best way to solve this problem is to augment the data, shift and scale the training data to match the holdout and or submission set and run it through a convolutional network. Using all the 40,000 training data, augmenting it and using a standard convolutional network will score 98%

But what if that wasn't an option? Using only 4000 labeled samples of 40,000 in the set and 4000 unlabeled samples accuracy can hit 81%

Using a Semi-Supervised GAN gives 5% better accuracy than the conventional method if you only have a limited amount of labeled data

## Kannada MNIST Semi-Supervised GAN

- Model small set of Kannada MNIST samples of 40,000 samples
- Holdout and Submission imgs vary from test data in std, mean

While the GAN slightly outperforms a standard Supervised FF Conv on Test Data that differs from Training Data it's not by a large amount. Augmenting data and shifting the std and mean to match Test Data is probably a better option

### Use labeled and unlabeled data from the training set ( same std, mean )

n_train samples	GAN validation	GAN holdout	Supervised validation	Supervised holdout
1000	98%	72%	97%	67%
2000	98%	70%	97%	62%
3000	98%	72%	98%	70%
4000	98%	72%	97%	67%
5000	99%	71%	98%	65%
6000	99%	76%	98%	67%
8000	99%	72%	98%	67%
10000	99%	75%	98%	70%
12000	99%	75%	99%	74%

### Use unlabeled holdout set as unlabeled data for GAN ( different std, mean) check against unlabeled set

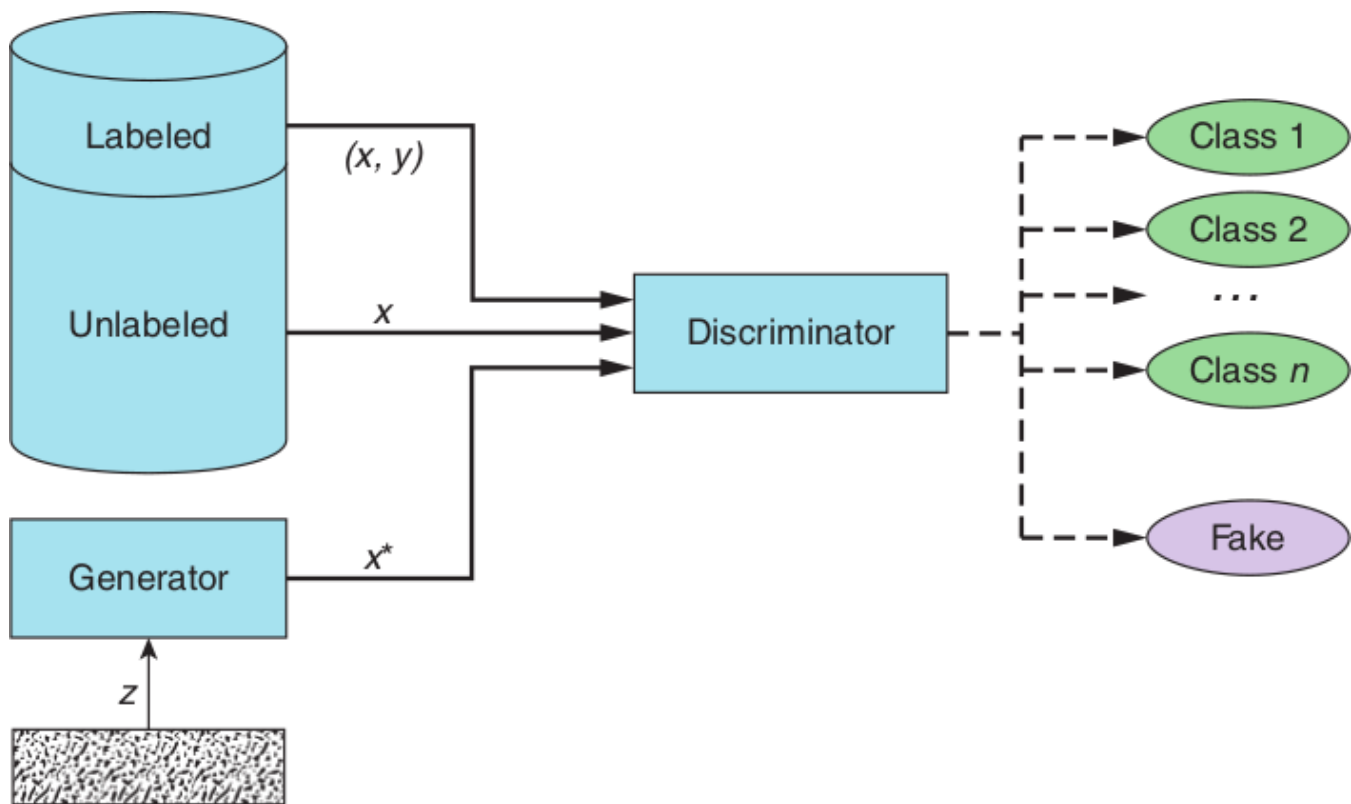
n_train samples	GAN validation	GAN holdout
1000	98%	79%
2000	98%	78%
3000	98%	%
4000	98%	80%
5000	99%	%
6000	99%	81%
8000	99%	%
10000	99%	72%
12000	99%	%

### Use Submission as unlabeled, check against holdout labeled data (all sets have diff std, mean)

4000 training samples yeilds 73% accuracy

- example code and images from <https://www.manning.com/books/gans-in-action> (<https://www.manning.com/books/gans-in-action>)

- dataset <https://www.kaggle.com/c/Kannada-MNIST> (<https://www.kaggle.com/c/Kannada-MNIST>)



- Reducing the Need for Labeled Data in Generative Adversarial Networks <https://ai.googleblog.com/2019/03/reducing-need-for-labeled-data-in.html> (<https://ai.googleblog.com/2019/03/reducing-need-for-labeled-data-in.html>)
- Self-Supervised GANs via Auxiliary Rotation Loss <https://arxiv.org/pdf/1811.11212.pdf> (<https://arxiv.org/pdf/1811.11212.pdf>)
- Unsupervised Representation Learning by Predicting Image Rotations <https://arxiv.org/abs/1803.07728> (<https://arxiv.org/abs/1803.07728>)

Kannada is a language spoken predominantly by people of Karnataka in southwestern India. The language has roughly 45 million native speakers and is written using the Kannada script.

The goal of this Kaggle competition is to use machine learning to correctly label hand-written digits written in the Kannada script. The Kannada dataset format is based on the dataset used in the original MNIST dataset where the objective was the same but the hand-written digits were Arabic numbers.

```

In [1]: 1 # silence is golden
        2
        3 import warnings
        4 warnings.filterwarnings("ignore")
        5 warnings.filterwarnings(action="ignore",category=DeprecationWarning)
        6 warnings.filterwarnings(action="ignore",category=FutureWarning)
        7
  
```

```

In [2]: 1 # hack to make keras work with 2*** series gpus
        2
        3 import tensorflow as tf
        4 config = tf.ConfigProto()
        5 config.gpu_options.allow_growth = True
        6 sess = tf.Session(config=config)
  
```

```
In [3]: 1 %matplotlib inline
2
3 import matplotlib.pyplot as plt
4 import numpy as np
5 import pandas as pd
6
7
8 from keras import backend as K
9
10 from keras.layers import (Activation, BatchNormalization, Concatenate, Dense
11                             Dropout, Flatten, Input, Lambda, Reshape)
12
13 from keras.layers.advanced_activations import LeakyReLU
14 from keras.layers.convolutional import Conv2D, Conv2DTranspose
15 from keras.models import Model, Sequential
16 from keras.optimizers import Adam
17 from keras.utils import to_categorical, print_summary
```

Using TensorFlow backend.

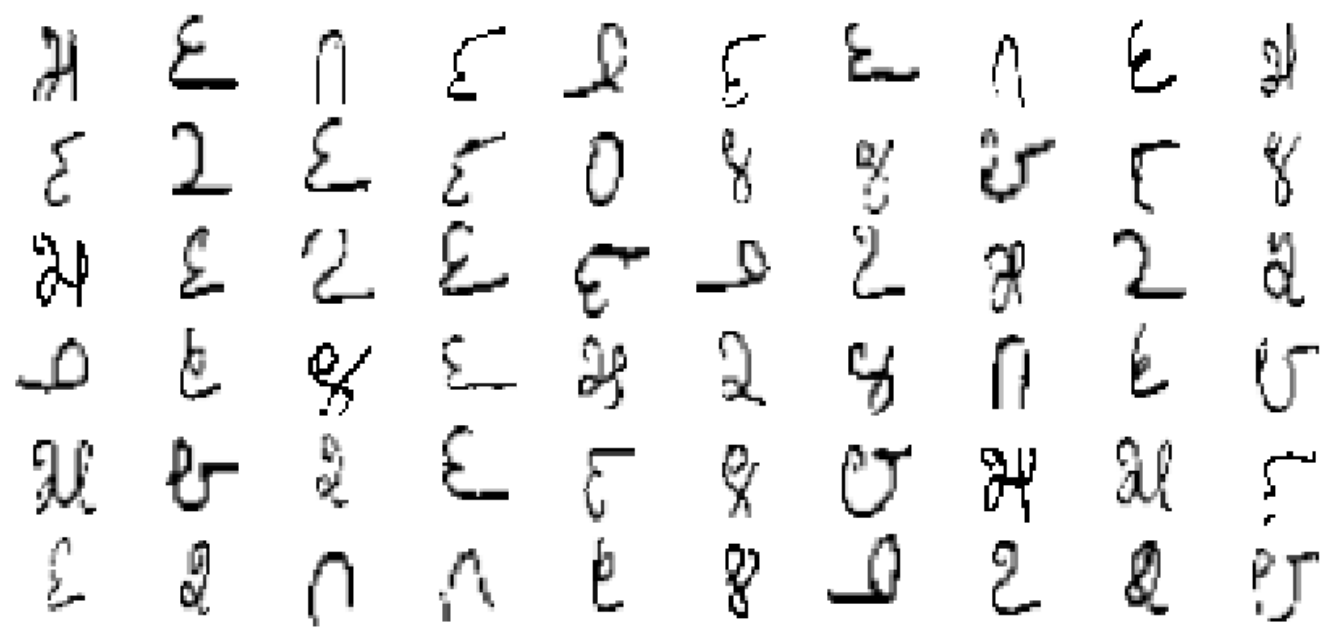
```
In [4]: 1 # Kannada MNIST dataset ~ 60,000 samples
2 # 40,000 labeled train
3 # 5,000 unlabeled submission
4 # 10,240 labeled holdout
5 # find labeled data, split out 100 samples for training
6
7
8 train = pd.read_csv('input/train.csv')
9 test = pd.read_csv('input/test.csv')
10 holdout = pd.read_csv('input/Dig-MNIST.csv')
11
12
13 # label, image
14 print('train', train.shape)
15 print(train.columns)
16
17
18 # id, image
19 print('test', test.shape)
20 print(test.columns)
21
22
23 # label, image
24 print('holdout', holdout.shape)
25 print(holdout.columns)
26
```

```
train (60000, 785)
Index(['label', 'pixel0', 'pixel1', 'pixel2', 'pixel3', 'pixel4', 'pixel5',
      'pixel6', 'pixel7', 'pixel8',
      ...,
      'pixel774', 'pixel775', 'pixel776', 'pixel777', 'pixel778', 'pixel779',
      'pixel780', 'pixel781', 'pixel782', 'pixel783'],
      dtype='object', length=785)
test (5000, 785)
Index(['id', 'pixel0', 'pixel1', 'pixel2', 'pixel3', 'pixel4', 'pixel5',
      'pixel6', 'pixel7', 'pixel8',
      ...,
      'pixel774', 'pixel775', 'pixel776', 'pixel777', 'pixel778', 'pixel779',
      'pixel780', 'pixel781', 'pixel782', 'pixel783'],
      dtype='object', length=785)
holdout (10240, 785)
Index(['label', 'pixel0', 'pixel1', 'pixel2', 'pixel3', 'pixel4', 'pixel5',
      'pixel6', 'pixel7', 'pixel8',
      ...,
      'pixel774', 'pixel775', 'pixel776', 'pixel777', 'pixel778', 'pixel779',
      'pixel780', 'pixel781', 'pixel782', 'pixel783'],
      dtype='object', length=785)
```

In [5]:

```
1
2 def show_imgs(df, n):
3
4     fig, ax = plt.subplots(n, 10)
5
6     for i in range(n):
7         for j in range(10):
8             r = np.random.randint(0, len(df))
9             number = df.iloc[r].values.reshape((28,28))
10            ax[i][j].imshow(number, cmap=plt.cm.binary)
11            ax[i][j].axis('off')
12
13        plt.subplots_adjust(wspace=0, hspace=0)
14        fig.set_figwidth(15)
15        fig.set_figheight(7)
16        fig.show()
17
18
19
20 print('Train')
21 train_imgs = train.drop(columns=['label'])
22 show_imgs(train_imgs, 6)
23
24 print('Test')
25 test_imgs = test.drop(columns=['id'])
26 show_imgs(test_imgs, 6)
27
28 print('Holdout')
29 holdout_imgs = holdout.drop(columns=['label'])
30 show_imgs(holdout_imgs, 6)
31
32
33
34
```

Train  
Test  
Holdout



၆	၇	၈	၉	၁၀	၁၁	၁၂	၁၃	၁၄	၁၅
၁၆	၁၇	၁၈	၁၉	၂၀	၂၁	၂၂	၂၃	၂၄	၂၅
၂၆	၂၇	၂၈	၂၉	၃၀	၃၁	၃၂	၃၃	၃၄	၃၅
၃၆	၃၇	၃၈	၃၉	၄၀	၄၁	၄၂	၄၃	၄၄	၄၅
၄၆	၄၇	၄၈	၄၉	၅၀	၅၁	၅၂	၅၃	၅၄	၅၅

၅၆	၅၇	၅၈	၅၉	၆၀	၆၁	၆၂	၆၃	၆၄	၆၅
၆၆	၆၇	၆၈	၆၉	၇၀	၇၁	၇၂	၇၃	၇၄	၇၅
၇၆	၇၇	၇၈	၇၉	၈၀	၈၁	၈၂	၈၃	၈၄	၈၅
၈၆	၈၇	၈၈	၈၉	၉၀	၉၁	၉၂	၉၃	၉၄	၉၅
၉၆	၉၇	၉၈	၉၉	၁၀၀	၁၀၁	၁၀၂	၁၀၃	၁၀၄	၁၀၅

In [6]:

```
1 # split into x image, y label, train -> train/validate
2
3 # this is the labeled part of the training set, rest of training data is held out
4 num_labeled = 4000
5
6
7 # shuffle train set
8 train_s = train.sample(frac=1.)
9 train = train_s[0:30000]
10 validate = train_s[30000:40000]
11
12 print(train.shape, validate.shape)
13
14
15 # split into labeled and unlabeled
16 train_labeled = train[0:num_labeled]
17 train_unlabeled = train[num_labeled:len(train)]
18
19
20 y_train_labeled = train_labeled.label.values
21 x_train_labeled = train_labeled.drop(columns=['label']).values
22
23 y_train_unlabeled = train_unlabeled.label.values
24 x_train_unlabeled = train_unlabeled.drop(columns=['label']).values
25
26
27 y_validate = validate.label.values
28 x_validate = validate.drop(columns=['label']).values
29
30
31
32 x_submission = test.drop(columns=['id']).values
33
34 y_holdout = holdout.label.values
35 x_holdout = holdout.drop(columns=['label']).values
36
37 print('labeled train', x_train_labeled.shape, y_train_labeled.shape)
38 print('unlabeled train', x_train_unlabeled.shape, y_train_unlabeled.shape)
39
40 print('submission', x_submission.shape)
41 print('holdout', x_holdout.shape, y_holdout.shape)
42
```

```
(30000, 785) (10000, 785)
labeled train (4000, 784) (4000,)
unlabeled train (26000, 784) (26000,)
submission (5000, 784)
holdout (10240, 784) (10240,)
```

```
In [7]: 1 # shift and scale image data
2 def scale_shift(x):
3     return x / 255.
4     #return (x-127.5)/(2 * x.std())
5
6
7
8 x_train_unlabeled = scale_shift(x_train_unlabeled)
9 x_train_labeled = scale_shift(x_train_labeled)
10 x_validate = scale_shift(x_validate)
11 x_holdout = scale_shift(x_holdout)
12 x_submission = scale_shift(x_submission)
13
14
15
16 print(x_train_unlabeled.std(), x_train_labeled.std(), x_validate.std(), x_holdout.std(), x_submission.std())
17 print(x_train_unlabeled.mean(), x_train_labeled.mean(), x_validate.mean(), x_holdout.mean(), x_submission.mean())
18 print(x_train_unlabeled.min(), x_train_labeled.min(), x_validate.min(), x_holdout.min(), x_submission.min())
19 print(x_train_unlabeled.max(), x_train_labeled.max(), x_validate.max(), x_holdout.max(), x_submission.max())
20
21
22
```

0.2418371564987552 0.24257205266890997 0.24149437347731595 0.2860085202946692 0.22543569156630952  
0.08234349585988249 0.08244536314525808 0.08211117997198893 0.11264889842655808 0.07261260104041618  
0.0 0.0 0.0 0.0 0.0  
1.0 1.0 1.0 1.0 1.0

```
In [8]: 1
2
3 # Expand image dimensions to width x height x channels
4 def reshape_img_data(x):
5     return np.expand_dims(x, axis=3)
6
7
8 print(x_train_unlabeled.shape, x_train_labeled.shape, x_validate.shape, x_holdout.shape, x_submission.shape)
9
10
11
12 # labels
13 def reshape_labels(y):
14     return y.reshape(-1, 1)
15
16 y_train_unlabeled = reshape_labels(y_train_unlabeled)
17 y_train_labeled = reshape_labels(y_train_labeled)
18 y_validate = reshape_labels(y_validate)
19 y_holdout = reshape_labels(y_holdout)
20
21 print(y_train_unlabeled.shape, y_train_labeled.shape, y_validate.shape, y_holdout.shape)
```

(26000, 784) (4000, 784) (10000, 784) (10240, 784) (5000, 784)  
(26000, 1) (4000, 1) (10000, 1) (10240, 1)

```
In [9]: 1 # train labels 0..9, equal numbers of each
2 print(y_train_unlabeled.min(), y_train_unlabeled.max())
3 print(y_train_labeled.min(), y_train_labeled.max())
4
5 print(np.unique(y_train_labeled, return_counts=True))
6 print(np.unique(y_train_unlabeled, return_counts=True))
7
8
9 print(y_holdout.min(), y_holdout.max())
10 print(np.unique(y_holdout, return_counts=True))
```

0 9  
0 9  
(array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9]), array([409, 444, 396, 409, 365, 384, 389, 398, 401, 405]))  
(array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9]), array([2555, 2591, 2521, 2612, 2675, 2556, 2563, 2657, 2603, 2667]))  
0 9  
(array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9]), array([1024, 1024, 1024, 1024, 1024, 1024, 1024, 1024, 1024, 1024]))

In [10]:

```
1 img_rows = 28
2 img_cols = 28
3 channels = 1
4
5 # Input image dimensions
6 img_shape = (img_rows, img_cols, channels)
7
8 # Size of the noise vector, used as input to the Generator
9 z_dim = 100
10
11 # Number of classes in the dataset
12 num_classes = 10
```



In [11]:

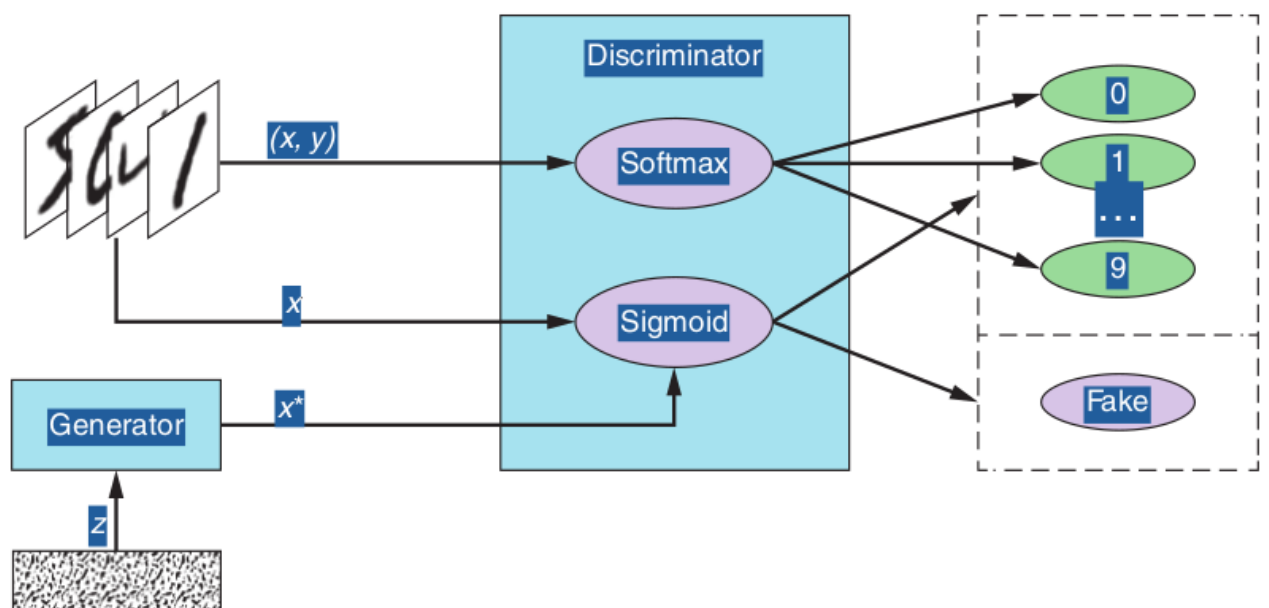
```
1
2
3
4
5 # random batch of labeled train data
6 def batch_labeled(batch_size):
7
8     x = x_train_labeled
9     y = y_train_labeled
10
11     # Get a random batch of labeled images and their labels
12     idx = np.random.randint(0, num_labeled, batch_size)
13     imgs = x[idx]
14     labels = y[idx]
15
16     imgs = imgs.reshape(batch_size, img_rows, img_cols, 1)
17     labels = to_categorical(labels, num_classes=num_classes)
18
19     return imgs, labels
20
21
22 # random batch unlabeled train data
23 def batch_unlabeled(batch_size):
24
25     x = x_train_unlabeled
26     y = y_train_unlabeled
27
28
29     # Get a random batch of unlabeled images
30     idx = np.random.randint(num_labeled, x.shape[0], batch_size)
31     imgs = x[idx]
32
33     imgs = imgs.reshape(batch_size, img_rows, img_cols, 1)
34
35
36     return imgs
37
38
39 # full training set
40 def train_set():
41
42     imgs = np.concatenate((x_train_labeled, x_train_unlabeled), axis=0)
43     y = np.concatenate((y_train_labeled, y_train_unlabeled), axis=0)
44
45     imgs = imgs.reshape(len(imgs), img_rows, img_cols, 1)
46     labels = to_categorical(y, num_classes=num_classes)
47
48
49     return imgs, labels
50
51
52
53
54 # validation set is a random selection of samples from train set
55 def test_set():
56
57     imgs = x_validate
58     y = y_validate
59
60     imgs = imgs.reshape(len(imgs), img_rows, img_cols, 1)
61     labels = to_categorical(y, num_classes=num_classes)
62
63     return imgs, labels
64
65
66
67 # unseen set, somewhat different std and mean in images
68 def holdout_set():
69
70     imgs = x_holdout
71     y = y_holdout
72
73     imgs = imgs.reshape(len(imgs), img_rows, img_cols, 1)
74     labels = to_categorical(y, num_classes=num_classes)
75
76     return imgs, labels
77
78
```

```

79 # unseen set, somewhat different std and mean in images
80 def submission_set():
81
82     imgs = x_submission
83     imgs = imgs.reshape(len(imgs), img_rows, img_cols, 1)
84
85     return imgs
86
87
88
89
90 #####
91 # experimental
92 # use unlabeled holdout and validation data as unlabeled instead of part
93 #   of the training set
94
95 # random batch unlabeled train data
96 def batch_holdout(batch_size):
97
98     x = x_holdout
99
100     # Get a random batch of unlabeled images
101     idx = np.random.randint(num_labeled, x.shape[0], batch_size)
102     imgs = x[idx]
103
104     imgs = imgs.reshape(batch_size, img_rows, img_cols, 1)
105
106
107     return imgs
108
109
110
111 # random batch unlabeled train data
112 def batch_submission(batch_size):
113
114     x = x_submission
115
116
117     # Get a random batch of unlabeled images
118     idx = np.random.randint(num_labeled, x.shape[0], batch_size)
119     imgs = x[idx]
120
121     imgs = imgs.reshape(batch_size, img_rows, img_cols, 1)
122
123
124     return imgs
125

```

## Semi-Supervised GAN



**Generator**

```
In [12]: 1 # model takes in a random noise vector of z_dim and learns to output a digit
2 # that is good enough that discriminator can't tell it from an MNIST database
3
4
5
6 def build_generator(z_dim):
7
8     model = Sequential()
9
10    # Reshape input into 7x7x256 tensor via a fully connected layer
11    model.add(Dense(256 * 7 * 7, input_dim=z_dim))
12    model.add(Reshape((7, 7, 256)))
13
14
15    # Transposed convolution layer, from 7x7x256 into 14x14x128 tensor
16    model.add(Conv2DTranspose(128, kernel_size=3, strides=2, padding='same'))
17    model.add(BatchNormalization())
18    model.add(LeakyReLU(alpha=0.01))
19
20
21    # Transposed convolution layer, from 14x14x128 to 14x14x64 tensor
22    model.add(Conv2DTranspose(64, kernel_size=3, strides=1, padding='same'))
23    model.add(BatchNormalization())
24    model.add(LeakyReLU(alpha=0.01))
25
26
27    # Transposed convolution layer, from 14x14x64 to 28x28x1 tensor
28    model.add(Conv2DTranspose(1, kernel_size=3, strides=2, padding='same'))
29    model.add(Activation('tanh'))
30
31
32    return model
```

## Discriminator

In [13]:

```
1  # takes input image and classifies it from 0...9
2
3
4  def build_discriminator_net(img_shape):
5
6      model = Sequential()
7
8      # Convolutional layer, from 28x28x1 into 14x14x32 tensor
9      model.add(
10         Conv2D(32,
11                kernel_size=3,
12                strides=2,
13                input_shape=img_shape,
14                padding='same'))
15      model.add(LeakyReLU(alpha=0.01))
16
17
18      # Convolutional layer, from 14x14x32 into 7x7x64 tensor
19      model.add(
20         Conv2D(64,
21                kernel_size=3,
22                strides=2,
23                input_shape=img_shape,
24                padding='same'))
25      model.add(BatchNormalization())
26      model.add(LeakyReLU(alpha=0.01))
27
28
29      # Convolutional layer, from 7x7x64 tensor into 3x3x128 tensor
30      model.add(
31         Conv2D(128,
32                kernel_size=3,
33                strides=2,
34                input_shape=img_shape,
35                padding='same'))
36      model.add(BatchNormalization())
37      model.add(LeakyReLU(alpha=0.01))
38      model.add(Dropout(0.5))
39
40
41      # Flatten the tensor
42      model.add(Flatten())
43      model.add(Dense(num_classes))
44
45      return model
```

In [14]:

```
1  # discriminator to label class
2  def build_discriminator_supervised(discriminator_net):
3
4      model = Sequential()
5      model.add(discriminator_net)
6      model.add(Activation('softmax'))
7
8      return model
```

In [15]:

```
1  # discriminator to id real/fake image
2  def build_discriminator_unsupervised(discriminator_net):
3
4      model = Sequential()
5      model.add(discriminator_net)
6
7      # Transform distribution over real classes into a binary real-vs-fake pr
8      def predict(x):
9          # 1 - 1/Sum(e^x + 1) --- almost sigmoid function
10         prediction = 1.0 - (1.0 / (K.sum(K.exp(x), axis=-1, keepdims=True) +
11         return prediction
12
13      # 'Real-vs-fake' output neuron defined above
14      model.add(Lambda(predict))
15
16      return model
```

## Build the Model

In [16]:

```
1 def build_gan(generator, discriminator):
2
3     model = Sequential()
4
5     # Combined Generator -> Discriminator model
6     model.add(generator)
7     model.add(discriminator)
8
9     return model
```

## Discriminator

```
In [17]: 1 # Core Discriminator network:
2 # These layers are shared during supervised and unsupervised training
3 discriminator_net = build_discriminator_net(img_shape)
4
5
6 # Build & compile the Discriminator for supervised training
7 discriminator_supervised = build_discriminator_supervised(discriminator_net)
8 discriminator_supervised.compile(loss='categorical_crossentropy', metrics=['accuracy'])
9 print('Supervised discriminator')
10 print_summary(discriminator_supervised)
11
12 # Build & compile the Discriminator for unsupervised training
13 discriminator_unsupervised = build_discriminator_unsupervised(discriminator_net)
14 discriminator_unsupervised.compile(loss='binary_crossentropy', optimizer=Adam())
15 print('Unsupervised discriminator')
16 print_summary(discriminator_unsupervised)
17
18
19
```

WARNING:tensorflow:From /home/herself/anaconda3/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:74: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

WARNING:tensorflow:From /home/herself/anaconda3/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:517: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /home/herself/anaconda3/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:4138: The name tf.random\_uniform is deprecated. Please use tf.random.uniform instead.

WARNING:tensorflow:From /home/herself/anaconda3/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:174: The name tf.get\_default\_session is deprecated. Please use tf.compat.v1.get\_default\_session instead.

WARNING:tensorflow:From /home/herself/anaconda3/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:181: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

WARNING:tensorflow:From /home/herself/anaconda3/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:1834: The name tf.nn.fused\_batch\_norm is deprecated. Please use tf.compat.v1.nn.fused\_batch\_norm instead.

WARNING:tensorflow:From /home/herself/anaconda3/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:3445: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.  
Instructions for updating:  
Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.  
WARNING:tensorflow:From /home/herself/anaconda3/lib/python3.7/site-packages/keras/optimizers.py:790: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

Supervised discriminator

Layer (type)	Output Shape	Param #
=====		
sequential_1 (Sequential)	(None, 10)	113930
-----		
activation_1 (Activation)	(None, 10)	0
=====		
Total params: 113,930		
Trainable params: 113,546		
Non-trainable params: 384		

WARNING:tensorflow:From /home/herself/anaconda3/lib/python3.7/site-packages/tensorflow/python/ops/nn\_impl.py:180: add\_dispatch\_support.<locals>.wrapper (from tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.  
Instructions for updating:  
Use tf.where in 2.0, which has the same broadcast rule as np.where  
Unsupervised discriminator

Layer (type)	Output Shape	Param #
=====		
sequential_1 (Sequential)	(None, 10)	113930

lambda_1 (Lambda)	(None, 1)	0
=====		
Total params: 113,930		
Trainable params: 113,546		
Non-trainable params: 384		
=====		

## Generator

```
In [18]: 1 # Build the Generator
2 generator = build_generator(z_dim)
3 print('Generator')
4 print_summary(generator)
5
6 # Keep Discriminator's parameters constant for Generator training
7 discriminator_unsupervised.trainable = False
8
9 # Build and compile GAN model with fixed Discriminator to train the Generator
10 # Note that we are using the Discriminator version with unsupervised output
11 gan = build_gan(generator, discriminator_unsupervised)
12 gan.compile(loss='binary_crossentropy', optimizer=Adam())
13 print('GAN')
14 print_summary(gan)
```

### Generator

Layer (type)	Output Shape	Param #
=====		
dense_2 (Dense)	(None, 12544)	1266944
reshape_1 (Reshape)	(None, 7, 7, 256)	0
conv2d_transpose_1 (Conv2DTr	(None, 14, 14, 128)	295040
batch_normalization_3 (Batch	(None, 14, 14, 128)	512
leaky_re_lu_4 (LeakyReLU)	(None, 14, 14, 128)	0
conv2d_transpose_2 (Conv2DTr	(None, 14, 14, 64)	73792
batch_normalization_4 (Batch	(None, 14, 14, 64)	256
leaky_re_lu_5 (LeakyReLU)	(None, 14, 14, 64)	0
conv2d_transpose_3 (Conv2DTr	(None, 28, 28, 1)	577
activation_2 (Activation)	(None, 28, 28, 1)	0
=====		
Total params: 1,637,121		
Trainable params: 1,636,737		
Non-trainable params: 384		

### GAN

Layer (type)	Output Shape	Param #
=====		
sequential_4 (Sequential)	(None, 28, 28, 1)	1637121
sequential_3 (Sequential)	(None, 1)	113930
=====		
Total params: 1,751,051		
Trainable params: 1,636,737		
Non-trainable params: 114,314		
=====		

## Training

In [19]:

```
1 supervised_losses = []
2 iteration_checkpoints = []
3
4
5
6 def train(iterations, batch_size, sample_interval):
7
8     # Labels for real images: all ones
9     real = np.ones((batch_size, 1))
10
11     # Labels for fake images: all zeros
12     fake = np.zeros((batch_size, 1))
13
14
15     for iteration in range(iterations):
16
17         # -----
18         # Train the Discriminator
19         # -----
20
21         # Get labeled examples
22         imgs, labels = batch_labeled(batch_size)
23
24         #####
25         # Get unlabeled examples
26         # imgs_unlabeled = batch_unlabeled(batch_size)
27         # imgs_unlabeled = batch_holdout(batch_size)
28         imgs_unlabeled = batch_submission(batch_size)
29
30         # Generate a batch of fake images
31         z = np.random.normal(0, 1, (batch_size, z_dim))
32         gen_imgs = generator.predict(z)
33
34         # Train on real labeled examples
35         d_loss_supervised, accuracy = discriminator_supervised.train_on_batch(
36
37         # Train on real unlabeled examples
38         d_loss_real = discriminator_unsupervised.train_on_batch( imgs_unlabeled, labels)
39
40         # Train on fake examples
41         d_loss_fake = discriminator_unsupervised.train_on_batch(gen_imgs, fake_labels)
42
43         d_loss_unsupervised = 0.5 * np.add(d_loss_real, d_loss_fake)
44
45         # -----
46         # Train the Generator
47         # -----
48
49         # Generate a batch of fake images
50         z = np.random.normal(0, 1, (batch_size, z_dim))
51         gen_imgs = generator.predict(z)
52
53         # Train Generator
54         g_loss = gan.train_on_batch(z, np.ones((batch_size, 1)))
55
56         if (iteration + 1) % sample_interval == 0:
57
58             # Save Discriminator supervised classification loss to be plotted
59             supervised_losses.append(d_loss_supervised)
60             iteration_checkpoints.append(iteration + 1)
61
62             # Output training progress
63             print(
64                 "%d [D loss supervised: %.4f, acc.: %.2f%] [D loss unsupervised: %.4f, acc.: %.2f%]"
65                 % (iteration + 1, d_loss_supervised, 100 * accuracy, d_loss_unsupervised, g_loss))
66
```

## Train the Model and Inspect Output

Note that the 'Discrepancy between trainable weights and collected trainable' warning from Keras is expected. It is by design: The Generator's trainable parameters are intentionally held constant during Discriminator training, and vice versa.

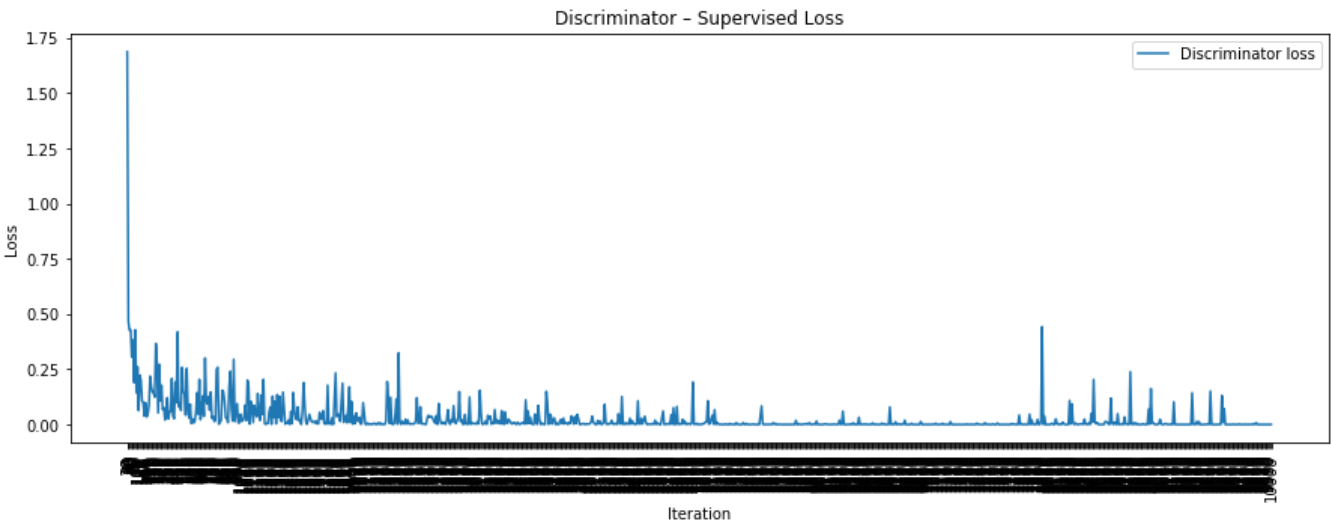


```
In [20]: 1 # Set hyperparameters
2 n_epochs = 10000
3 batch_size = 32
4 sample_interval = 10
5
6 # Train the SGAN for the specified number of iterations
7 train(n_epochs, batch_size, sample_interval)

10 [D loss supervised: 1.6872, acc.: 46.88%] [D loss unsupervised: 0.2600] [G
loss: 1.899957]
20 [D loss supervised: 0.4659, acc.: 90.62%] [D loss unsupervised: 0.0328] [G
loss: 0.474590]
30 [D loss supervised: 0.4263, acc.: 90.62%] [D loss unsupervised: 0.0345] [G
loss: 0.861463]
40 [D loss supervised: 0.4284, acc.: 81.25%] [D loss unsupervised: 0.0162] [G
loss: 0.874039]
50 [D loss supervised: 0.3062, acc.: 93.75%] [D loss unsupervised: 0.0058] [G
loss: 0.199279]
60 [D loss supervised: 0.3830, acc.: 87.50%] [D loss unsupervised: 0.0086] [G
loss: 0.147347]
70 [D loss supervised: 0.1902, acc.: 96.88%] [D loss unsupervised: 0.0053] [G
loss: 0.253196]
80 [D loss supervised: 0.4286, acc.: 84.38%] [D loss unsupervised: 0.0074] [G
loss: 0.331915]
90 [D loss supervised: 0.1436, acc.: 93.75%] [D loss unsupervised: 0.0044] [G
loss: 0.023412]
100 [D loss supervised: 0.2604, acc.: 90.62%] [D loss unsupervised: 0.0028] [G
loss: 0.012043]
```

```
In [21]: 1 losses = np.array(supervised_losses)
2
3 # Plot Discriminator supervised loss
4 plt.figure(figsize=(15, 5))
5 plt.plot(iteration_checkpoints, losses, label="Discriminator loss")
6
7 plt.xticks(iteration_checkpoints, rotation=90)
8
9 plt.title("Discriminator – Supervised Loss")
10 plt.xlabel("Iteration")
11 plt.ylabel("Loss")
12 plt.legend()
```

Out[21]: <matplotlib.legend.Legend at 0x7f1c58490c50>



## SGAN Classifier – Training, Test, Holdout Accuracy

```
In [22]: 1 # compute training data accuaracy
2 x, y = train_set()
3
4 _, accuracy = discriminator_supervised.evaluate(x, y)
5 print("Test Accuracy: %.2f%%" % (100 * accuracy))

30000/30000 [=====] - 1s 37us/step
Test Accuracy: 98.47%
```

In [23]:

```
1 # compute validation data accuracy
2 x, y = test_set()
3
4 _, accuracy = discriminator_supervised.evaluate(x, y)
5 print("Test Accuracy: %.2f%%" % (100 * accuracy))
```

10000/10000 [=====] - 0s 36us/step

Test Accuracy: 98.41%

In [24]:

```
1 # compute holdout data accuracy
2 x, y = holdout_set()
3
4 _, accuracy = discriminator_supervised.evaluate(x, y)
5 print("Test Accuracy: %.2f%%" % (100 * accuracy))
```

10240/10240 [=====] - 0s 40us/step

Test Accuracy: 73.44%

In [25]:

```
1 # create Kaggle submission data
2 # score was not good enough to bother submitting, code here only to show how
3 x = submission_set()
4
5 gan_predictions = discriminator_supervised.predict(x)
6 print('gan preds', gan_predictions.shape)
7 gan_preds = np.argmax(gan_predictions, axis = 1)
8
9 print(gan_preds)
```

gan preds (5000, 10)

[3 0 2 ... 1 6 3]

## Fully-Supervised Classifier

- use same number of labeled samples
- see diff in number of training epoches needed

In [26]:

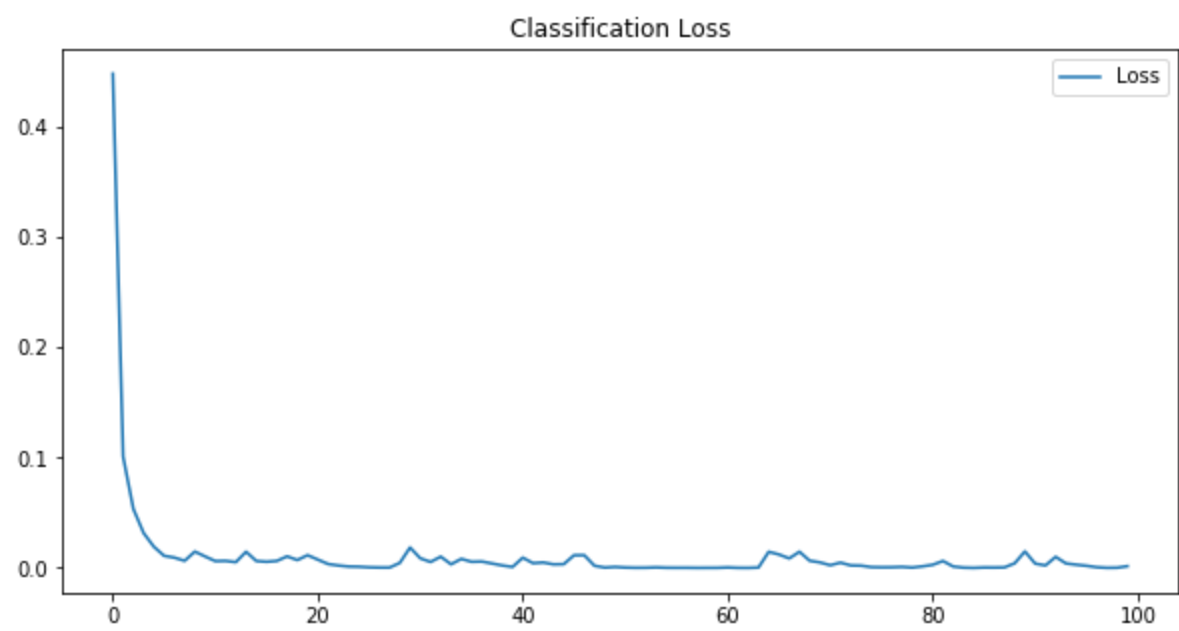
```
1 # Fully supervised classifier with the same network architecture as the SGAN
2 mnist_classifier = build_discriminator_supervised(build_discriminator_net(im
3 mnist_classifier.compile(loss='categorical_crossentropy', metrics=['accuracy
```

```
In [27]: 1 imgs, labels = batch_labeled(num_labeled)
2
3
4 # Train the classifier
5 training = mnist_classifier.fit(x=imgs,
6                                y=labels,
7                                batch_size=batch_size,
8                                epochs=n_epochs//100,
9                                verbose=1)
10
11 losses = training.history['loss']
12 accuracies = training.history['acc']
```

Epoch 1/100  
4000/4000 [=====] - 1s 249us/step - loss: 0.4484 - acc: 0.8560  
Epoch 2/100  
4000/4000 [=====] - 0s 107us/step - loss: 0.1011 - acc: 0.9667  
Epoch 3/100  
4000/4000 [=====] - 0s 105us/step - loss: 0.0536 - acc: 0.9848  
Epoch 4/100  
4000/4000 [=====] - 0s 109us/step - loss: 0.0317 - acc: 0.9933  
Epoch 5/100  
4000/4000 [=====] - 0s 114us/step - loss: 0.0192 - acc: 0.9955  
Epoch 6/100  
4000/4000 [=====] - 0s 105us/step - loss: 0.0110 - acc: 0.9978  
Epoch 7/100  
4000/4000 [=====] - 0s 101us/step - loss: 0.0084 - acc: 0.9994

```
In [28]: 1 # Plot classification loss
2 plt.figure(figsize=(10, 5))
3 plt.plot(np.array(losses), label="Loss")
4 plt.title("Classification Loss")
5 plt.legend()
```

Out[28]: <matplotlib.legend.Legend at 0x7f1bf0d07128>



```
In [29]: 1 # check training data accuracy
2
3 x, y = train_set()
4 _, accuracy = mnist_classifier.evaluate(x, y)
5 print("Training Accuracy: %.2f%%" % (100 * accuracy))
```

30000/30000 [=====] - 1s 40us/step  
Training Accuracy: 97.95%

In [30]:

```
1 # check validation data accuracy
2
3 x, y = test_set()
4 _, accuracy = mnist_classifier.evaluate(x, y)
5 print("Test Accuracy: %.2f%%" % (100 * accuracy))
```

10000/10000 [=====] - 0s 38us/step  
Test Accuracy: 97.96%

In [31]:

```
1 # check holdout data accuracy
2
3 x, y = holdout_set()
4 _, accuracy = mnist_classifier.evaluate(x, y)
5 print("Test Accuracy: %.2f%%" % (100 * accuracy))
```

10240/10240 [=====] - 0s 40us/step  
Test Accuracy: 71.66%

In [32]:

```
1 # create Kaggle submission data
2
3 x = submission_set()
4
5 clf_predictions = mnist_classifier.predict(x)
6 print('clf preds', clf_predictions.shape)
7 clf_preds = np.argmax(clf_predictions, axis = 1)
8 print(clf_preds)
```

clf preds (5000, 10)  
[3 0 2 ... 1 6 3]

---