

Numerical Computation and Machine Learning Basics

<https://github.com/timestocome/DeepLearning-Talks>

Overflow and Underflow

Overflow -+/- infinity, NaN, wrapping

Underflow - numbers fade to zero

vanishing gradients [clip gradients, add epsilon]

exploding gradients [clip gradients, use regularization]

- Pentium FDIV Bug
- GPU truncation causing non-linearity
- <https://blog.openai.com/nonlinear-computation-in-linear-networks/>

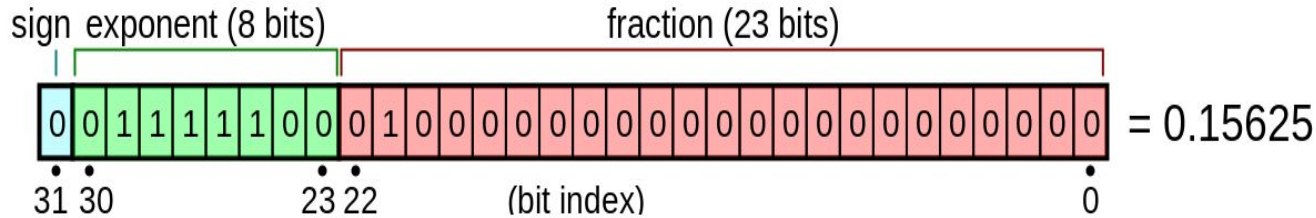
Most ML libraries will handle NaN, -/+ inf, div by zero for you, or add epsilon (a very small number) to denominator

Wrapping: use a 16 bit int to count in a loop and it'll wrap at 65,536 - 1

- don't use loops
- use a very large loop index

IEEE Standard

$(-1 * s) \times (\text{fraction}) \times (2 ^ \text{exponent})$



(* fraction is usually named mantissa)

Size of ints, longs, floats, doubles on your computer

<https://www.programiz.com/c-programming/examples/sizeof-operator-example>

http://cstl-csm.semo.edu/xzhang/Class%20Folder/CS280/Workbook_HTML/FLOATING_tut.htm

http://www.boundedfloatingpoint.com/PressRelease_011718.pdf (end of floating point errors ?)

Poor Conditioning

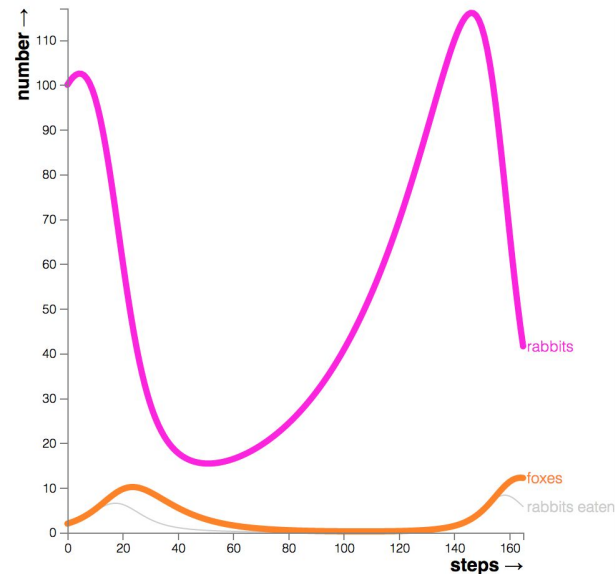
Conditioning is how quickly a function changes

Even a small learning rate increases the error causing the network to become stuck

In a small network this is a local minima, in a deep network it's a saddle point

Fix: use an adaptive learning rate (Adagrad, AdaDelta, RMSProp, Adam, Momentum ..

Adaptive learning rates are built into Theano, TensorFlow ...

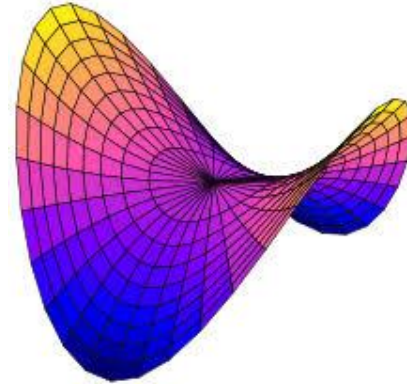
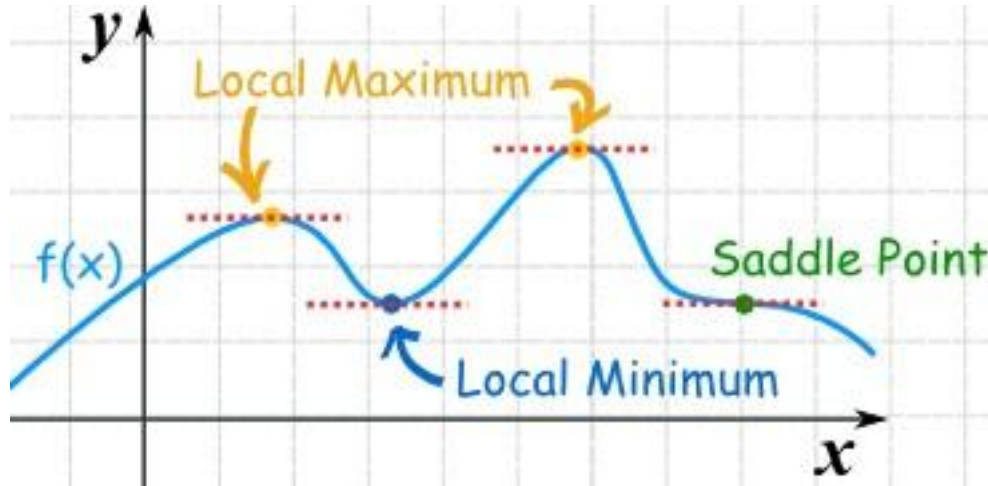


Gradient Based Optimization

Goal is to minimize or maximize a cost function (loss function, error function)

Using a small learning rate, take a small step in the direction of steepest decrease

Gradient = 0 at minimum, max, saddle point



Derivative


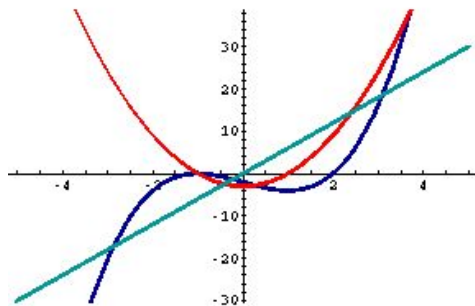
Measures the tangent line on a slope which gives you the direction of greatest increase. In neural networks we use the negative of it to get steepest descent.

The 2nd derivative can tell you if you are at a max or min

$f'' < 0$ max

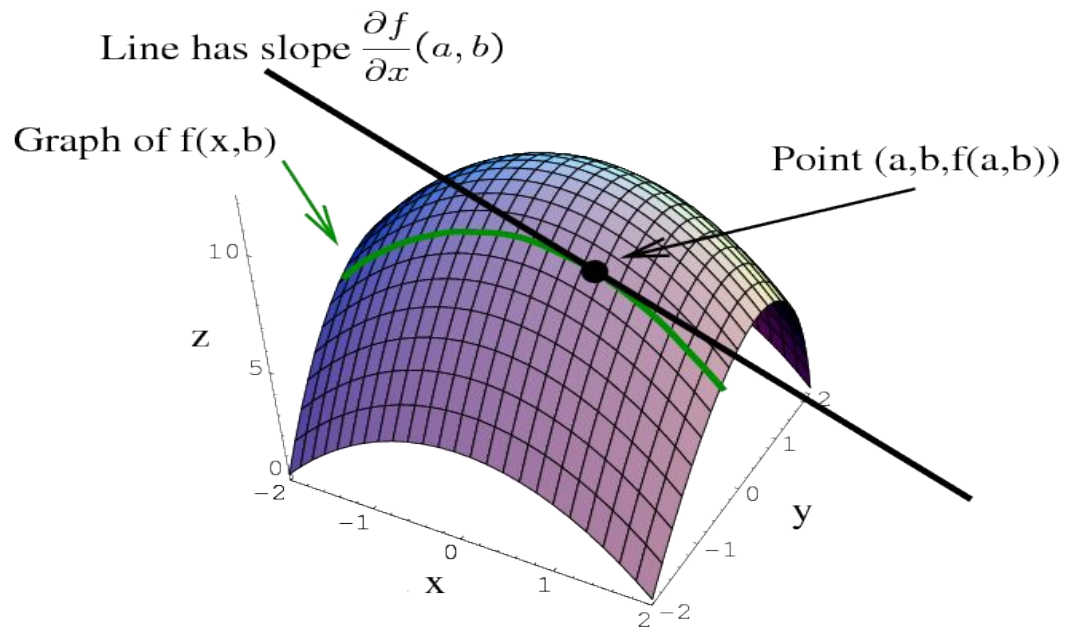
$f'' = 0$ no information

$f'' > 0$ minimum



Partial Derivative


Derivatives calculated on vectors (gradient)



Jacobian Matrix

Take the first derivative of a vector and output a vector to find steepest descent

$$\frac{\partial (x, y)}{\partial (u, v)} = \begin{vmatrix} \frac{\partial x}{\partial u} & \frac{\partial x}{\partial v} \\ \frac{\partial y}{\partial u} & \frac{\partial y}{\partial v} \end{vmatrix} = \frac{\partial x}{\partial u} \frac{\partial y}{\partial v} - \frac{\partial x}{\partial v} \frac{\partial y}{\partial u}$$



Hessian Matrix

Used to find the 2nd derivative of a vector

Measures the curvature

Use to adjust learning rate (step size)

negative -> maxes negative curvature is down,

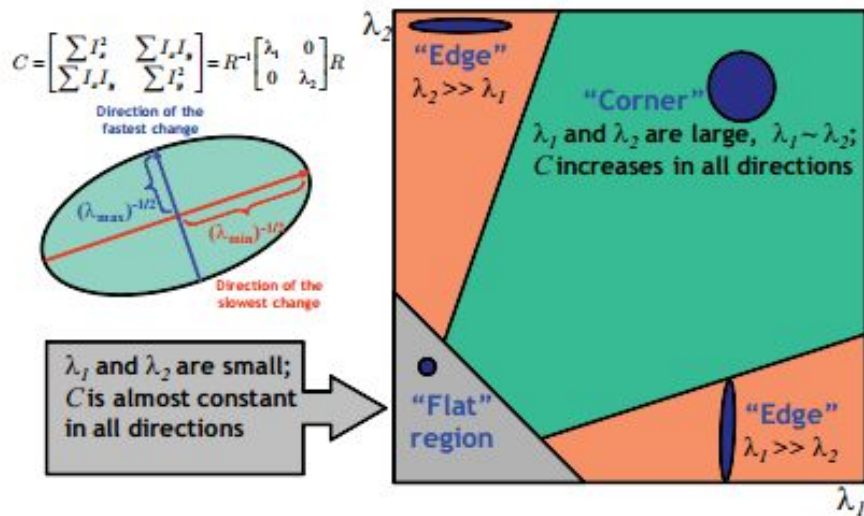
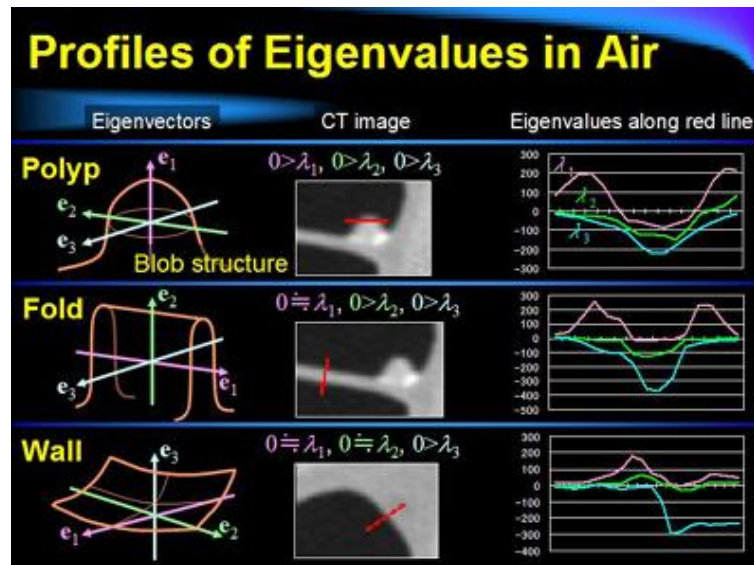
positive -> mins curvature is up

zero -> flat space or a saddle point

$$\text{Gradient} = \begin{bmatrix} \frac{\delta H}{\delta z_1} \\ \frac{\delta H}{\delta z_2} \\ \frac{\delta H}{\delta z_3} \end{bmatrix}, \text{Hessian} = \begin{bmatrix} \frac{\partial^2 H}{\partial z_1^2} & \frac{\partial^2 H}{\partial z_1 \partial z_2} & \frac{\partial^2 H}{\partial z_1 \partial z_3} \\ \frac{\partial^2 H}{\partial z_2 \partial z_1} & \frac{\partial^2 H}{\partial z_2^2} & \frac{\partial^2 H}{\partial z_2 \partial z_3} \\ \frac{\partial^2 H}{\partial z_3 \partial z_1} & \frac{\partial^2 H}{\partial z_3 \partial z_2} & \frac{\partial^2 H}{\partial z_3^2} \end{bmatrix}$$

Eigenvalues of Hessian

Find ridges (fingerprints, CAT scans); Motion capture (character animation)



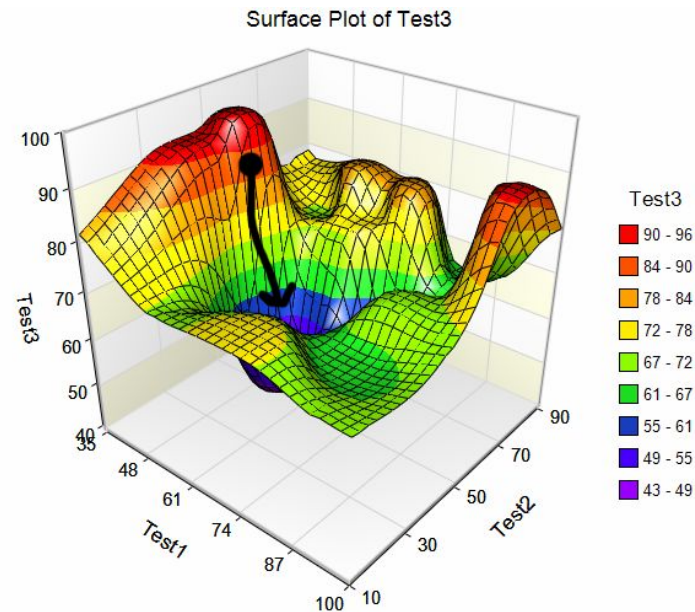
Stochastic Gradient Descent

Repeat until convergence

{

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$


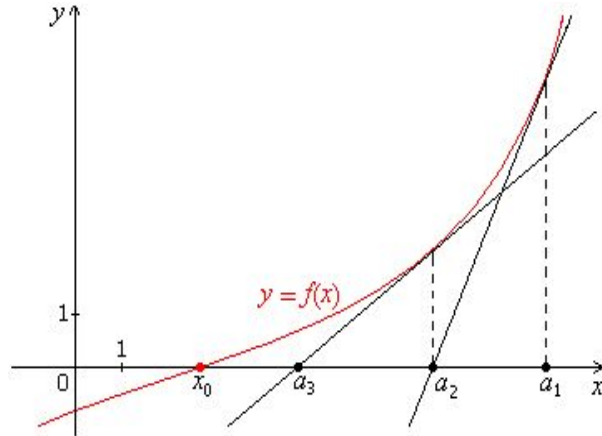
}



Newton's Method

Approximate $f(x) = 0$

Keep guessing, split the difference between guesses until $f(x) = 0$ Usually fails spectacularly, step size missed zero



Code example: <http://www.aip.de/groups/soe/local/numres/bookcpdf/c9-4.pdf>

Constrained Optimization

More than one cost in the problem

i.e. Find production number that maximizes income using only factory available

Or maximize $f(x)$ subject to $g(x)$

Often just add $g(x)$ to cost function as a penalty, but a large penalty can create ill conditioned Hessians causing the network to fail to converge

Lagrangian Multiplier

solve: $\max \log(x) + \log(y)$ given: $x + y = m$

Combine equations, set $g(x)$ to zero $\max \log(x) + \log(y) + L(m - x - y)$

Set derivatives = 0 $1/x - z = 0, 1/y - z = 0, m - x - y = 0$

Remove L from x, y $1/x - z = 1/y - z \implies x = y$

2 equations, 2 unknowns $x = y = m/2$

Solve for Lagrangian multiplier $z = m/2$

<http://home.uchicago.edu/~vlima/courses/econ201/pricetext/chapter2.pdf>

http://adl.stanford.edu/aa222/Lecture_Notes_files/constrainedOptimization.pdf

<https://www.khanacademy.org/math/multivariable-calculus/applications-of-multivariable-derivatives/constrained-optimization/a/lagrange-multipliers-single-constraint>

Karush-Kuhn-Tucker Approach

- Combining Lagrange conditions for equality and inequality constraints yields *KKT conditions* for general

problem: $\min_{\mathbf{x}} f(\mathbf{x})$

s. t. $\mathbf{g}(\mathbf{x}) \leq \mathbf{0}$

$\mathbf{h}(\mathbf{x}) = \mathbf{0}$

Lagrangian:

$$L = f(\mathbf{x}) + \boldsymbol{\mu}^T \mathbf{g}(\mathbf{x}) + \boldsymbol{\lambda}^T \mathbf{h}(\mathbf{x})$$

\Rightarrow

$$\frac{\partial L}{\partial \mathbf{x}} = \frac{\partial f}{\partial \mathbf{x}} + \sum \mu_i \frac{\partial g_i}{\partial \mathbf{x}} + \sum \lambda_i \frac{\partial h_i}{\partial \mathbf{x}} = \mathbf{0} \quad (\text{optimality})$$

and

$$\mathbf{g} \leq \mathbf{0}, \quad \mathbf{h} = \mathbf{0} \quad (\text{feasibility})$$

$$\boldsymbol{\lambda} \neq \mathbf{0}, \quad \boldsymbol{\mu} \geq \mathbf{0}, \quad \mu_i g_i = 0 \quad (\text{complementarity})$$

Linear Least Squares with Constraint

Solve: $f(x) = \frac{1}{2} \|Ax - b\|^2$

Constraint: $g(x) = x.Tx \leq 1$

Combine the equations: $\frac{1}{2}\|Ax-b\|^2 + z(x.Tx - 1)$

set derivative = 0 $A.TAx - A.Tb + 2zx = 0$

Solve for Lagrangian Multiplier $z = (A.TAx - A.Tb) * x/2$

and the constraint $dL/dz = x.Tx - 1$

Learning Algorithms

Able to learn from data

"A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance in tasks at tasks in T , as measured by P , improves with experience E ."

E : data, sensor input, images,

T : goals

P : error

ML is the field of study that gives computers the ability to learn without being explicitly programmed - Arthur Samuel

Experience (data, sensor inputs...)

Missing values

- Use most common value, average or median
- Interpolation between values (time series)
- Drop that feature
- Infer from other features

Categories Use one hot vectors

Sensors Allow for noise

Independent Identically Distributed

Garbage in, garbage out

Independent

- features cannot be correlated

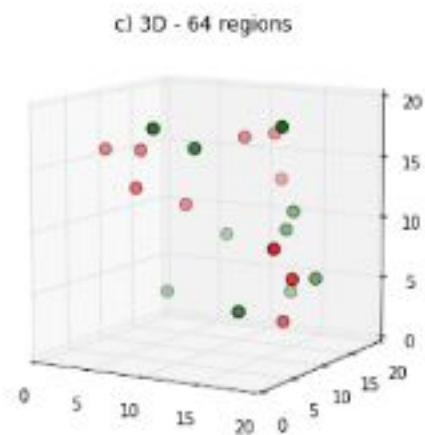
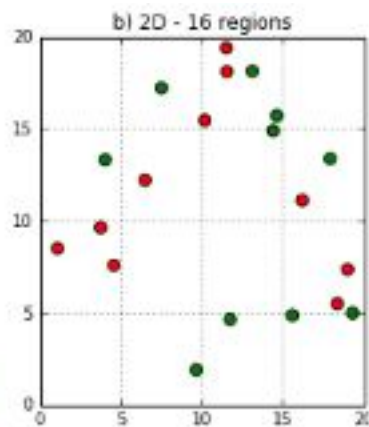
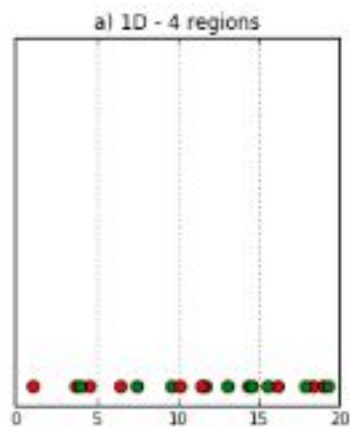
Identically Distributed

- training, validation, test data
- outputs (if 90% of the output is true, the network will always guess true and you'll have 90% accuracy - check the confusion matrix)

<https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.corr.html>

http://scikit-learn.org/stable/modules/generated/sklearn.metrics.confusion_matrix.html

Curse of dimensionality



The Task (Goals)

This is what you want your neural network to solve, the goal

Think carefully, you can get very different answers depending on how you frame your task.

-- First self driving cars were trained with human drivers and learned to drive on their own.

-- Test car is successfully driven down the road, then down a hill into a lake

The Performance Measure (Cost function)

~ 80% of data is used for training

~ 10% of data to validate

~ 10% of data is held out to be used as a test after training is done

While accuracy less than (pick a number)%:

- shuffle data

- run training data in batches

- run validation every 100th epoch or so:

Then test on hold out data after training is complete (* never let subcontractors, contest entrants see hold out data)

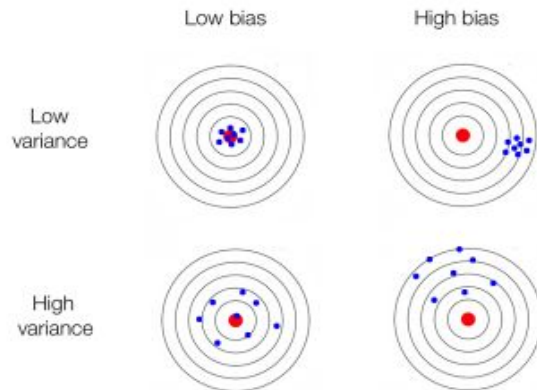
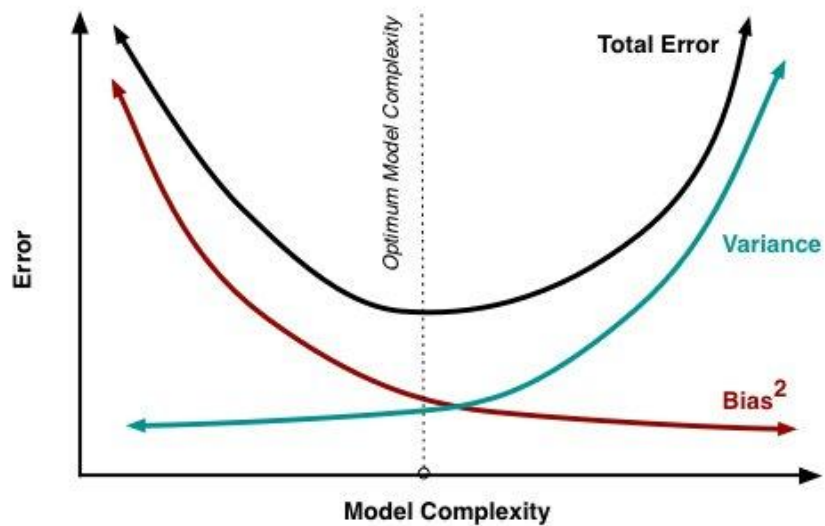
Training, testing, validation data must be statistically similar

Bias vs Variance

High bias == underfitting

High variance == overfitting

$$\text{MSE} = (\text{target} - \text{predicted})^2 = \text{Bias}^2 + \text{Variance} + \text{noise}$$



Entropy

Entropy = how many bits are needed to encode information

(* information theory uses Log2, not e, not 10)

Coin toss: 2 possibilities (heads, tails) == 1 bit

Entropy = - Sum($(\frac{1}{2}) * \log(\frac{1}{2}) + (\frac{1}{2}) * \log(\frac{1}{2})$)

Entropy = - ($\frac{1}{2} * -1 + \frac{1}{2} * -1$) = -(- $\frac{1}{2}$ - $\frac{1}{2}$) = 1

Coin toss: 2 possibilities (heads $\frac{1}{3}$, tails $\frac{2}{3}$)

Entropy = -($\frac{1}{3} * \log(\frac{1}{3}) + \frac{2}{3} * \log(\frac{2}{3})$) = 0.89

A Mathematical Theory of Communication, Shannon

<http://affect-reason-utility.com/1301/4/shannon1948.pdf>

$$\text{Entropy} = - \sum_i P_i \log_2 P_i$$

KL Divergence

Kullback-Leibler Divergence tells how much information is lost between the neural network prediction and the actual values

$$D_{KL}(p||q) = \sum_{i=1}^N p(x_i) \cdot (\log p(x_i) - \log q(x_i))$$

p is probability distribution (actual), q is approximation (predicted)

The principle of maximum likelihood says that given the training data, we should use as gives the greatest possible probability to the training data

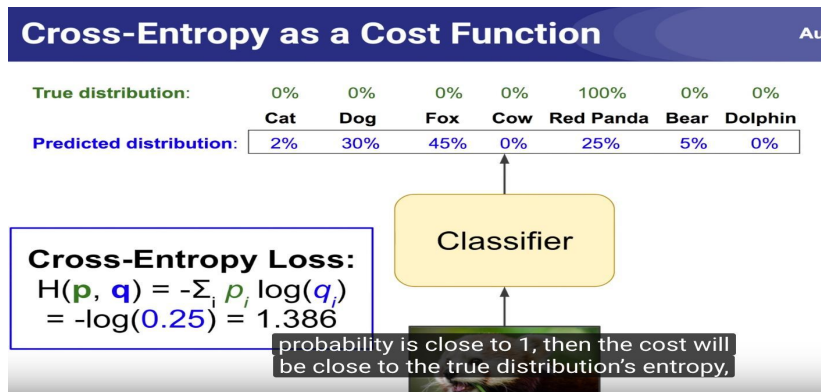
- not a distance measure !

<https://www.countbayesie.com/blog/2017/5/9/kullback-leibler-divergence-explained>

<http://cseweb.ucsd.edu/~elkan/250B/logreg.pdf>

$$\text{Entropy} = - \sum_i P_i \log_2 P_i$$

Gross entropy as cost function



Cross entropy = entropy + KL Divergence

entropy = minimum bits needed to transmit actual information

KL Divergence = difference between predicted and actual information

<https://www.youtube.com/watch?v=ErfnhcEV1O8&feature=youtu.be>

Bayesian vs Frequentist

Bayesian

Prior beliefs

Fixed parameters

$$\mu = mx + b$$

$$y = N(\mu, \text{std})$$

N - normal distribution

std - standard deviation

μ - mean


Frequentist

Repeatable random sample


Fixed data

$$y = mx + b$$


DID THE SUN JUST EXPLODE?
(IT'S NIGHT, SO WE'RE NOT SURE.)



FREQUENTIST STATISTICIAN:



BAYESIAN STATISTICIAN:



Bayesian Statistics

The Posterior	The Evidence	The Prior
	The probability of getting this evidence if this hypothesis were true	The probability of H being true, before gathering evidence
$P(H E) = \frac{P(H E)P(H)}{P(E)}$		
The probability that the hypothesis (H) is true given the evidence (E)	The marginal probability of the evidence (Prob of E over all possibilities)	

Intuitive understanding of Bayes

https://arbital.com/p/bayes_frequency_diagram/?l=55z&pathId=28771

http://scikit-learn.org/stable/modules/naive_bayes.html

Relative size	Case B	Case \bar{B}	Total
Condition A	w	x	w+x
Condition \bar{A}	y	z	y+z
Total	w+y	x+z	w+x+y+z

$$P(A|B) \times P(B) = \frac{w}{w+y} \times \frac{w+y}{w+x+y+z} = \frac{w}{w+x+y+z}$$

$$P(B|A) \times P(A) = \frac{w}{w+x} \times \frac{w+x}{w+x+y+z} = \frac{w}{w+x+y+z}$$

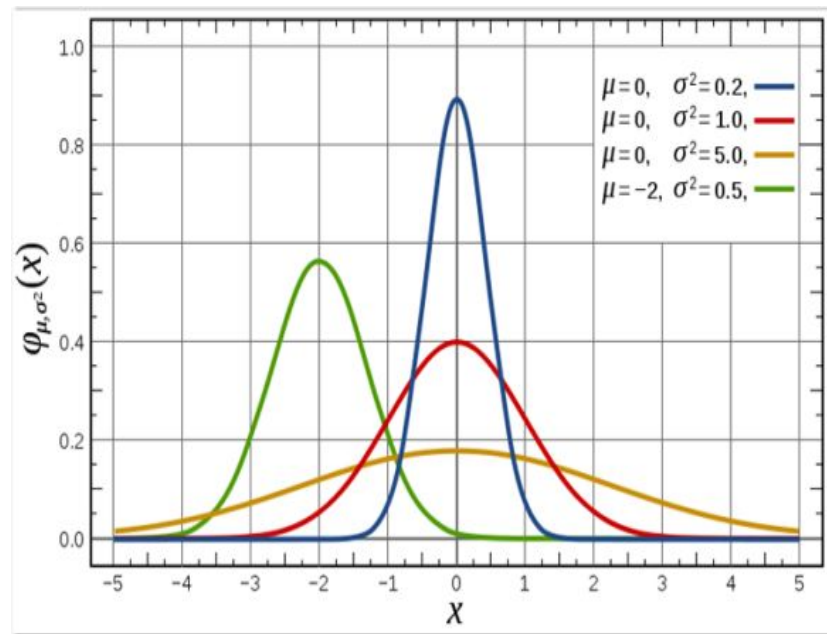
Maximum Log Likelihood (frequentist)

probability of observing the given data as a function of theta

$$\theta_{MLE} = \arg \max_{\theta} \log P(X|\theta)$$

$$= \arg \max_{\theta} \log \prod_i P(x_i|\theta)$$

$$= \arg \max_{\theta} \sum_i \log P(x_i|\theta)$$



Conditional Log Likelihood (Bayesian version)

Weights are trained to maximize the conditional data likelihood

$$W^* = \operatorname{argmax}_W \prod_{d \in D} P(Y^d \mid X_1^d \dots X_n^d)$$

which is the same as maximizing the conditional log likelihood

$$W^* = \operatorname{argmax}_W \sum_{d \in D} \ln P(Y^d \mid X_1^d \dots X_n^d)$$

<http://cseweb.ucsd.edu/~elkan/250B/logreg.pdf>

Mean Squared Error

If the data examples are i.i.d.

independent and identically distributed, then the Conditional Log Likelihood can be reduced to the Mean Squared Error

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \tilde{y}_i)^2$$

Linear Regression as Maximum Likelihood

Dot product as a measure of similarity between vectors

MSE \sim Conditional log likelihood

As $n_{\text{samples}} \rightarrow \infty$ convergence increases

- True distribution of data must be in model family
- True distribution must correspond to one set of weights

Let's take the Con out of Econometrics

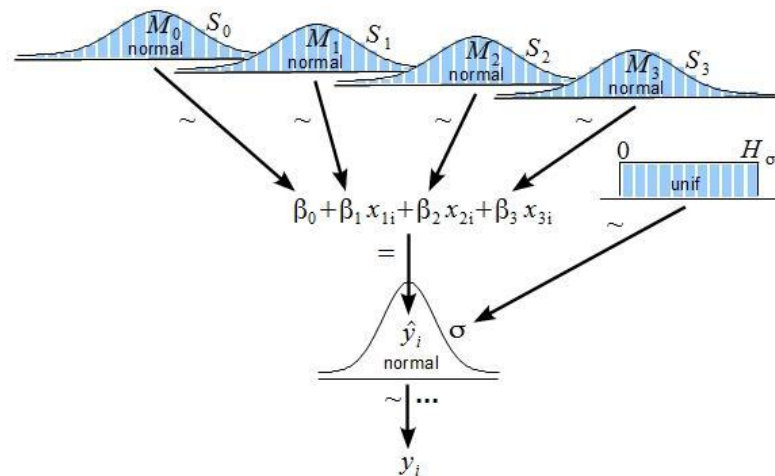
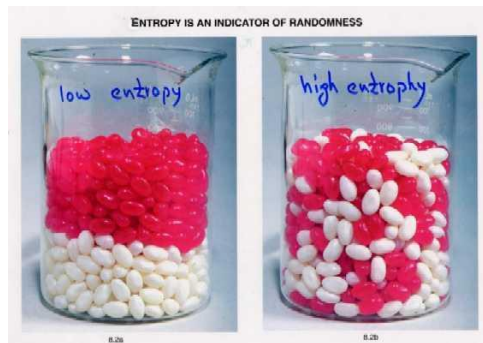
<http://www.econ.ucla.edu/workingpapers/wp239.pdf>

Bayesian Linear Regression

- Prior is uniform or Gaussian distribution with high entropy
- Observation reduces entropy and drives weights to single values

Using a known prior probability

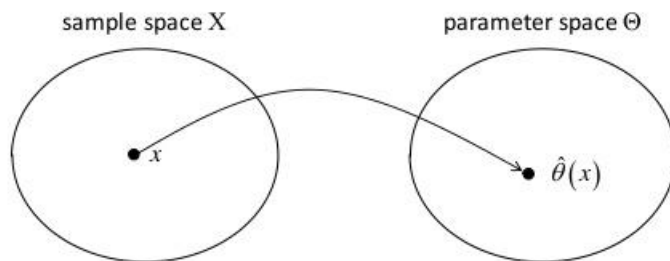
Prior distribution shifts weights toward simpler smoother curves



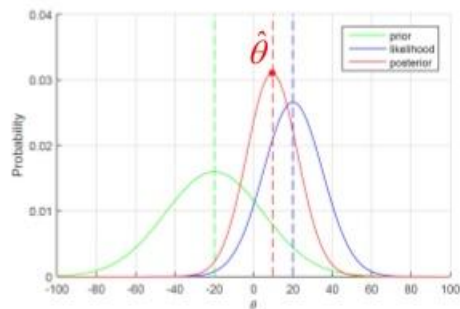
<https://stats.stackexchange.com/questions/252577/bayes-regression-how-is-it-done-in-comparison-to-standard-regression>

Maximum a Posteriori (MAP) Estimation

Bayesian estimation: Maximum A Posteriori (MAP).



$$\begin{aligned}\hat{\theta}(x) &= \arg \max_{\theta} P(\theta | x) \\ &= \arg \max_{\theta} \frac{P(x | \theta) P(\theta)}{P(x)} \\ &= \arg \max_{\theta} P(x | \theta) P(\theta)\end{aligned}$$



MLE vs MAP

Both compute a single value, not a distribution

MLE - fit a Gaussian to the data

MAP - uses posterior distribution

MLE is a specific case of MAP

$$\begin{aligned}\theta_{MAP} &= \arg \max_{\theta} \sum_i \log P(x_i|\theta)P(\theta) \\ &= \arg \max_{\theta} \sum_i \log P(x_i|\theta) \text{ const} \\ &= \arg \max_{\theta} \sum_i \log P(x_i|\theta) \\ &= \theta_{MLE}\end{aligned}$$

$$\begin{aligned}\theta_{MLE} &= \arg \max_{\theta} \log P(X|\theta) \\ &= \arg \max_{\theta} \log \prod_i P(x_i|\theta) \\ &= \arg \max_{\theta} \sum_i \log P(x_i|\theta)\end{aligned}$$

<https://wiseodd.github.io/techblog/2017/01/01/mle-vs-map/>

Cross validation

Split the dataset into 3 sections:

Train on one section

Test on one section

Validate on one section

Shuffle data, split and repeat

Average error gives you an approximate idea of how well an algorithm will perform

Especially useful on small datasets.

http://scikit-learn.org/stable/modules/cross_validation.html

Hyperparameters

Variables used to control the algorithm

- max, min, depth of branches on a decision tree
- number of layers, size of layers in network
- learning rate, complexity constant

Usually a grid search (nested loops over several options) is used with cross validation

http://scikit-learn.org/stable/modules/grid_search.html

A recent development is that randomly picking values for a grid search is more effective than stepping through all values.

<http://www.jmlr.org/papers/volume13/bergstra12a/bergstra12a.pdf>

Regularization

Training adjusts the weights to minimize the error, regularization forces the weights to stay small which forces the network to generalize

Used to prevent overfitting: To be useful a network must be able to generalize to new data.

- L1 regularization
- L2 regularization
- Early stopping
- Drop out
- Batch normalization

L1 Regularization

aka Lasso

Minimize the sum of the absolute differences

Drives some weights to zero which has the effect of pruning features which acts as a feature selector. Almost never used in real problems because it is not rotationally invariant

$L1 = \text{complexity_parameter} * \sum(\text{abs}(W))$

L2 Regularization

aka Ridge regression

Limits weight size without driving any to zero, minimizes variance, increases bias

Rotational invariance is the main reason it is used instead of L1

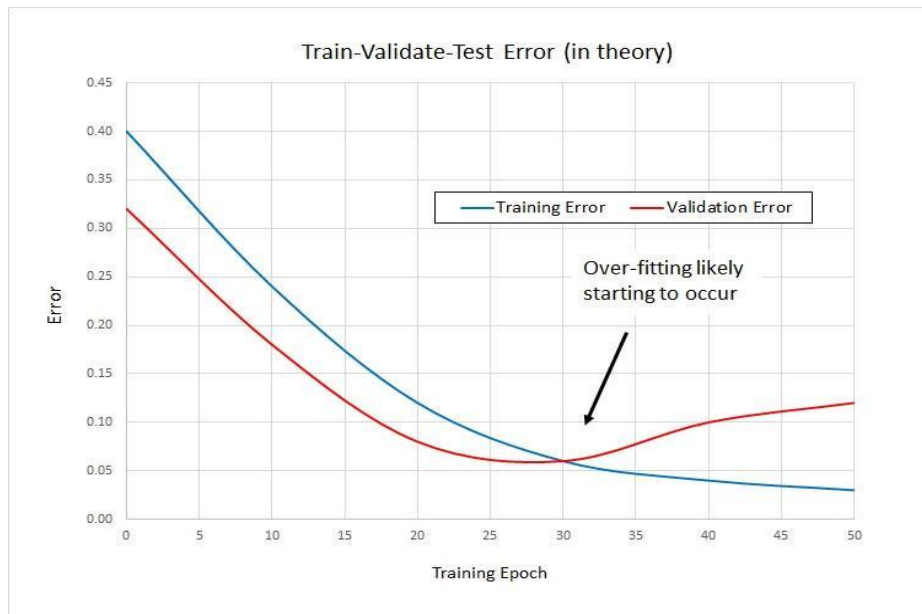
$$L2 = \text{complexity_constant} * \sum(W)^2$$

- L1 and L2 have the effect of choosing a different prior in Bayesian solutions
- Elastic net uses both L1 and L2

[https://web.stanford.edu/~hastie/Papers/B67.2%20\(2005\)%20301-320%20Zou%20&%20Hastie.pdf](https://web.stanford.edu/~hastie/Papers/B67.2%20(2005)%20301-320%20Zou%20&%20Hastie.pdf)

Capacity, Overfitting, Underfitting

Early stopping



		prediction outcome		total
		p	n	
actual value	p'	True Positive	False Negative	P'
	n'	False Positive	True Negative	N'
total		P	N	

Dropout

Averages predictions over all weights


Randomly exclude a percent of hidden nodes from the neural network, change and randomly exclude a different group on next batch during training.

How to:

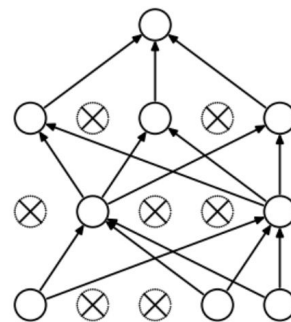
- Create a matrix the size of the weights matrix

- Randomly fill with ones and zeros

- Multiply weights by the matrix



(a) Standard Neural Net



(b) After applying dropout.

Dropout: A Simple Way to Prevent NN from Overfitting

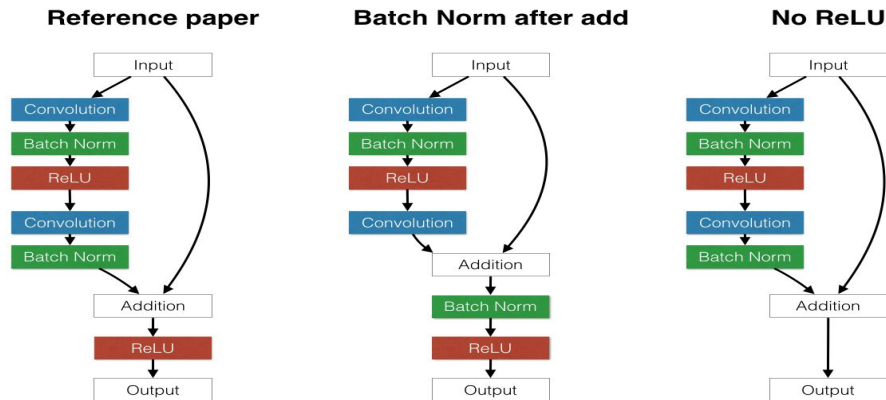
<https://www.cs.toronto.edu/~hinton/absps/JMLRdropout.pdf>

Batch Normalization

Is often used in place of DropOut, can use higher learning rates, networks train faster.

For each batch of training examples recenter the features.

In practice, the original input is fed to later layers



<https://arxiv.org/pdf/1502.03167.pdf>

Probabilistic Supervised Learning

Normal (Gaussian Distribution)

Probability of y given x

The activation functions are used to create non-linearity

- sigmoid, oldest, may cause vanishing gradients
- tanh, steeper than sigmoid, still can cause vanishing gradients
- relu, not linear, combinations or Relu also nonlinear, can approximate any function, drives some weights to zero
- fast to compute
- leaky relu (use when relu drives all your weights to zero)
-
- softmax (used in last layer to scale values so they sum up to 1)

Linear Regression

prediction = dot(w.T, x)

MSE = 1/n_samples * Sum(predicted - actual)^2

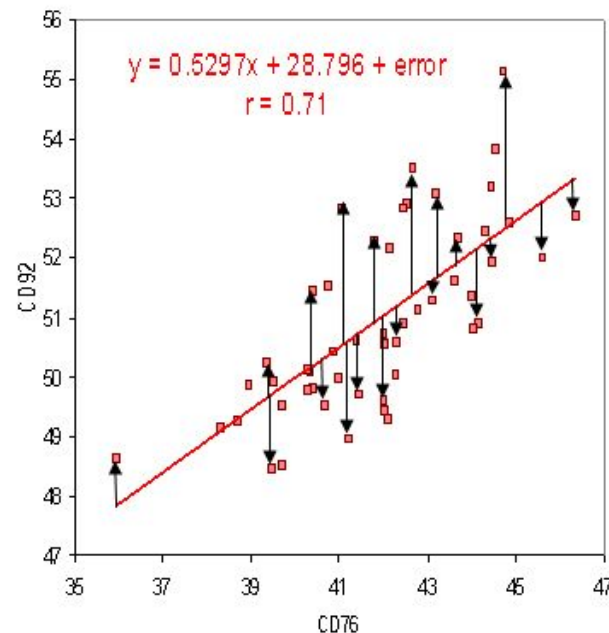

Solve for unknown $f(x)$

$J(w) = ||y - \text{dot}(w.T, x)||^2$

Use derivative to find direction to change weights

$dJ/dw = 2 * X^T X w - 2 * X^T y$

$w = w - \text{learning_rate} * 2 * (X^T X w - X^T y)$



Classification

Classification: image recognition, medical symptoms, marketing, voters

Features are turned into vectors and the distance between vectors is used to determine group membership. (city block, bird's eye)

Features representing values are scaled between 0,1 or -1,1

Features representing classes are converted to one hot vectors

Outputs a one hot vector selecting a class

Classification is unsupervised learning and might not always classify in a useful way

Different runs with data that's been shuffled between runs can classify the data differently

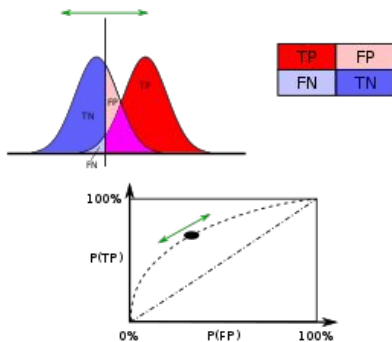
K-Means Clustering

Pick n points at random (centroids), must give algorithm the number of clusters

Put each point in the group of the closest centroid using either Manhattan or Euclidean distance, repeat until stable


- results may vary on sorting order

ROC Receiver operating characteristic ~ 1



<http://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html>

<http://stanford.edu/~cpiech/cs221/handouts/kmeans.html>



Classification with missing inputs

Joint Training of Deep Boltzmann Machines for Classification, Goodfellow (Goodfellow 2013b paper)

<https://arxiv.org/pdf/1301.3568.pdf>

Restricted Deep Boltzmann Machines, Salakhutdinov and Hinton

<http://proceedings.mlr.press/v5/salakhutdinov09a/salakhutdinov09a.pdf>

Deep Learning 4j has a nice tutorial on Boltzmann machines

<https://deeplearning4j.org/restrictedboltzmannmachine>

(aka Deep Belief Networks)



Transcription

OCR, Optical Character Recognition

Multi-digit Number Recognition from Street View

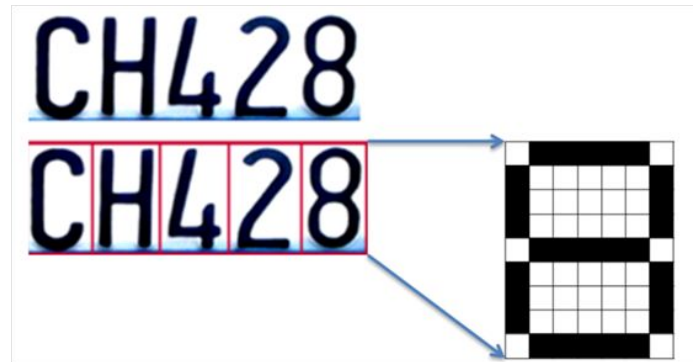
<https://static.googleusercontent.com/media/research.google.com/en//pubs/archive/42241.pdf>

Goodfellow talk on the paper

https://www.youtube.com/watch?v=vGPI_JvLoN0

Dataset

<http://ufldl.stanford.edu/housenumbers/>



Machine Translation

Used for natural language translation

Typically done with a recurrent network (RNN, LSTM, GRU, ...)

One language is the input, the second language is the output.

The network is trained back and forth rather than just forward feeding data and back feeding the errors

<https://nlp.stanford.edu/>

<https://github.com/MicrosoftTranslator/DocumentTranslator>

Structured Output

Collobert Home Page

<https://scholar.google.com/citations?user=32w7x1cAAAAJ&hl=en>

Natural Language Processing (Almost) from Scratch (Collobert 2012)

<http://www.jmlr.org/papers/volume12/collobert11a/collobert11a.pdf>

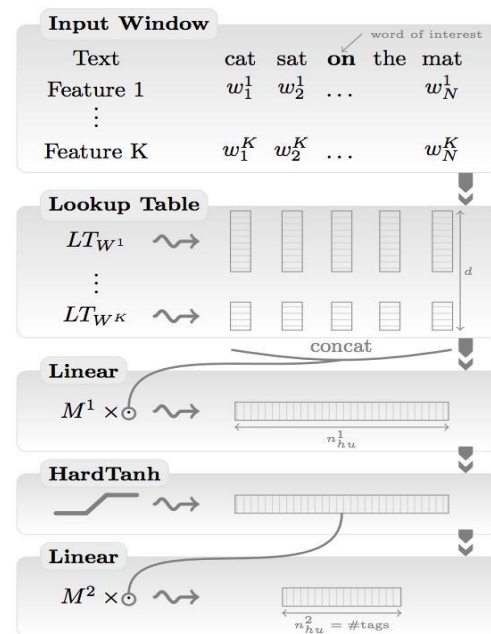


Figure 1: Window approach network.

Anomaly Detection

Make a prediction, points outside the prediction area are anomalies. This is a high value research area -- many problems still to be solved (equipment time to failure, stock market crashes, earthquake prediction.....)



RAD - Outlier Detection on Big Data

<https://medium.com/netflix-techblog/rad-outlier-detection-on-big-data-d6b0494371cc>

Bayesian Anomaly detection

<http://www.dtic.mil/cgi-bin/GetTRDoc?AD=ADA610860>

Support Vector Machines



$O(n^2)$, can only separate 2 classes per SVM

<http://scikit-learn.org/stable/modules/svm.html>

https://docs.opencv.org/2.4.13.4/doc/tutorials/ml/introduction_to_svm/introduction_to_svm.html

Decision Tree

Use entropy to make splits to give largest information gain

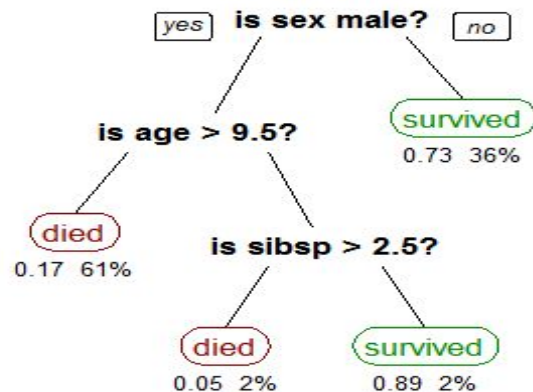
<http://scikit-learn.org/stable/modules/tree.html>

XGBoost

<http://xgboost.readthedocs.io/en/latest/model.html>

Using decision trees to see what a neural network is doing

<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.12.6011&rep=rep1&type=pdf>




Principal Components Analysis

Used to compress data to reduce features by combining them in a way that maximizes variance

Use this when data is highly correlated to create independent feature representations

<https://onlinecourses.science.psu.edu/stat857/node/35>

<http://scikit-learn.org/stable/modules/generated/sklearn.decomposition.PCA.html>



Manifold Learning

Find a useful way to represent the data in less dimensions

Example: represent road as 2D in a 3D world

It's like PCA but attempts to retain more information

t-SNE Used often for visualization

Doesn't always work, information can be lost or distorted

<http://scikit-learn.org/stable/modules/manifold.html>

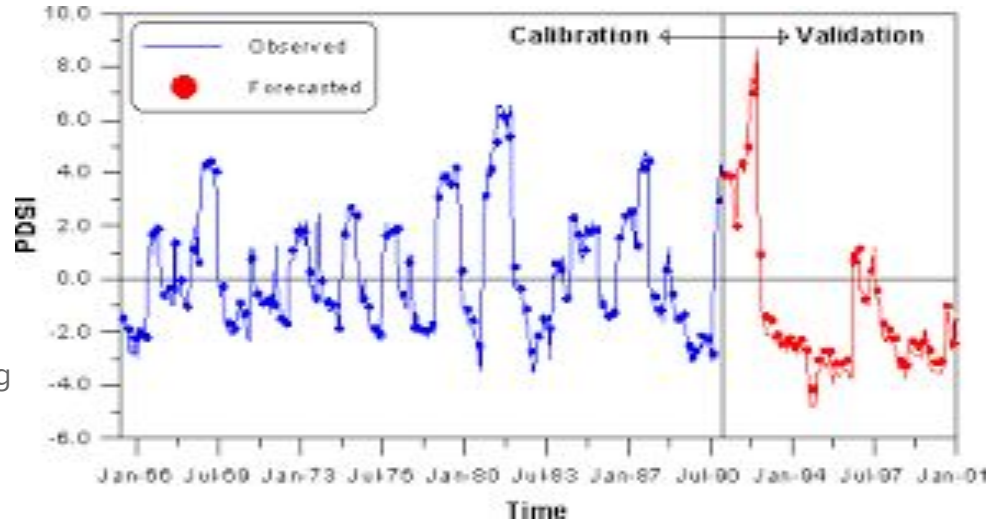
Time Series Predictions

Used to predict future values

- Must be stationary (rotate down to x axis)
- Generally use a log scale
- Validation data must be future data

Facebook's Prophet is a great open source tool for doing

<https://facebook.github.io/prophet/>



Synthesis and Sampling

Used in the game industry to create landscapes (fractal), characters (AI), etc

<https://www.gamasutra.com/topic/game-developer>

Microsoft Speech recognition and synthesis

<https://github.com/Microsoft/Windows-universal-samples/tree/master/Samples/SpeechRecognitionAndSynthesis>

<https://www.microsoft.com/en-us/research/blog/microsoft-researchers-achieve-new-conversational-speech-recognition-milestone/>

<https://arxiv.org/pdf/1708.06073.pdf>

No Free Lunch Theorem

Says that if you average over all the machine learning algorithms each will have the same performance on the same data set.

The only way one strategy can outperform another is by adjusting the network to the problem.

- convolutional networks for vision
- recursive networks for time series
- deep vs wide layers

<http://www.asimovinstitute.org/neural-network-zoo/>

Building a Machine Learning Algorithm

Cleanup data

Multiply data by weights, iterative training for non-linear data

Minimize difference between cost function and actual data

(* no free lunch theorem)

New old things being tried to adjust weights

NEAT, NeuroEvolution of Augmenting Topologies <http://nn.cs.utexas.edu/downloads/papers/stanley.ec02.pdf>

Genetic Algorithms https://www.doc.ic.ac.uk/~nd/surprise_96/journal/vol1/hmw/article1.html

Reinforcement Learning <http://www.incompleteideas.net/book/bookdraft2018jan1.pdf>

Videos

Bay Area Deep Learning School 2016

<https://www.youtube.com/watch?v=eyovmAtUx0> and <https://www.youtube.com/watch?v=9dXiAecyJrY>

Deep Learning Summer School Montreal

<https://www.youtube.com/playlist?list=PL5bqlc6XopCbb-FvnHmD1neVIQKwGzQyR>

http://videolectures.net/deeplearning2017_montreal/

Neural Networks for Machine Learning (Hinton Lectures 78 videos)

https://www.youtube.com/playlist?list=PLoRI3Ht4JOcdU872GhiYWf6jwrk_SNhz9

Blogs and things

ArXiv Papers <https://arxiv.org/>

Code for Arxiv papers: <http://www.gitxiv.com/>

Uber ML Blog: <https://eng.uber.com/>

Deep Mind <https://deepmind.com/>

Netflix ML Blog: <https://medium.com/@NetflixTechBlog>

Aylien NLP Blog and Data: <http://blog.aylien.com/research/>

Open AI: <https://openai.com/systems> Test environments for RL learning

Kaggle: <https://blog.kaggle.com/> Data, ML Blog, forums, examples, contests

Tensorflow

Tensorflow <https://www.tensorflow.org/>

Tensorflow for poets: <https://codelabs.developers.google.com/codelabs/tensorflow-for-poets/#0>

<https://www.youtube.com/watch?v=cSKfRcEDGUs>

Github: <https://github.com/tensorflow/tensorflow>

TensorFlow Models <https://github.com/tensorflow/models>

Manning MEAP Books <https://www.manning.com/meap-catalog> (*these assume you know ML and don't know TF)

Numerical computing

Numerical Recipes in C/C++, Fortran 77/90

<http://numerical.recipes/oldverswitcher.html>

Introduction to Statistical Learning

<http://www-bcf.usc.edu/~gareth/ISL/ISLR%20Seventh%20Printing.pdf>

Elements of Statistical Learning

<https://web.stanford.edu/~hastie/Papers/ESLII.pdf>

Deep Learning

Neural Networks and Deep Learning <http://neuralnetworksanddeeplearning.com/>

Must Know Tricks in Deep Neural Networks <http://lamda.nju.edu.cn/weixs/project/CNNTricks/CNNTricks.html>

Neural Network Zoo <http://www.asimovinstitute.org/neural-network-zoo/>

Neural Network Playground <http://playground.tensorflow.org>

WildML <http://www.wildml.com/>

Software

Anaconda <https://www.anaconda.com/>

Microsoft CNTK <https://github.com/Microsoft/CNTK>

Deep Learning for Java <https://deeplearning4j.org/index.html>

Deep Learn JS <https://deeplearnjs.org/>

Keras <https://keras.io/#keras-the-python-deep-learning-library>

XGBoost - Extreme Gradient Boosting <https://github.com/dmlc/xgboost>

Data: <https://archive.ics.uci.edu/ml/index.php>