Numerical Computation

Topics in Chapter 4 of Deep Learning

https://github.com/timestocome/DeepLearning-Talks

Numerical computing resources:

Numerical Recipes in C/C++, Fortran 77/90

http://numerical.recipes/oldverswitcher.html

Accelerated Numerical Analysis Tools with GPUs

https://developer.nvidia.com/how-to-numerical-analysis

Overflow and Underflow

Overflow -+/- infinity, NaN, wrapping

Underflow - numbers fade to zero, integers wrap
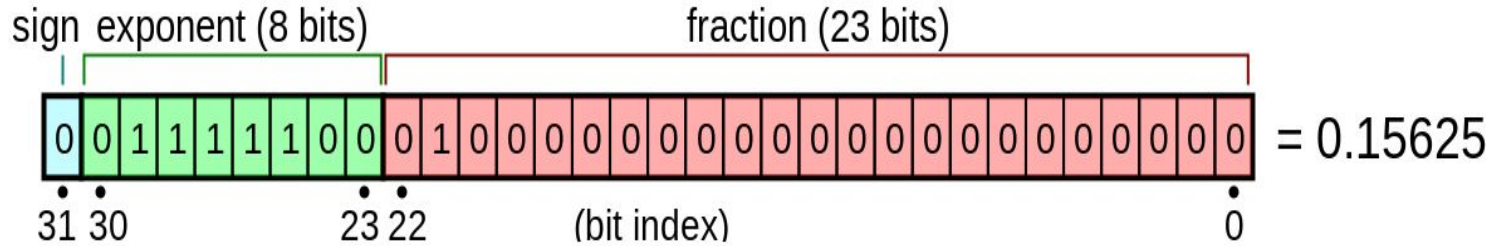
vanishing gradients [clip gradients, add epsilon]

exploding gradients [clip gradients, use regularization]

- Pentium FDIV Bug
- GPU truncation causing non-linearity
- https://blog.openai.com/nonlinear-computation-in-linear-networks/

Most ML libraries will handle NaN for you

Wrapping: use a 16 bit int to count in a loop and it'll wrap at 65,536 - 1

IEEE Standard

Size of ints, longs, floats, doubles on your computer

https://www.programiz.com/c-programming/examples/sizeof-operator-example

http://cstl-csm.semo.edu/xzhang/Class%20Folder/CS280/Workbook_HTML/FLOATING_tut.htm

Poor Conditioning

Even a small learning rate increases the error causing the network to become stuck

In a small network this is a local minima, in a deep network it's a saddle point

Fix: use an adaptive learning rate (Adagrad, AdaDelta, RMSProp, Adam, Momentum ...)

Adaptive learning rates are built into Theano, TensorFlow ...

Gradient Based Optimization

Goal is to minimize or maximize a cost function (loss function, error function)
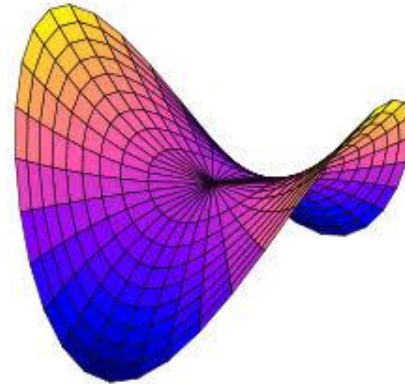
Derivative

Measures the tangent line on a slope which gives you the direction of greatest increase. In neural networks we use the negative of it to get steepest descent. The 2nd derivative can tell you if you are at a max or min
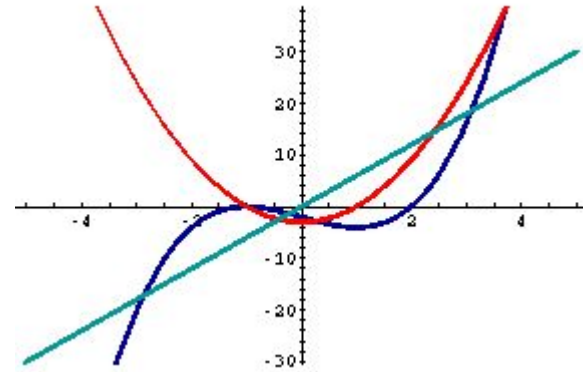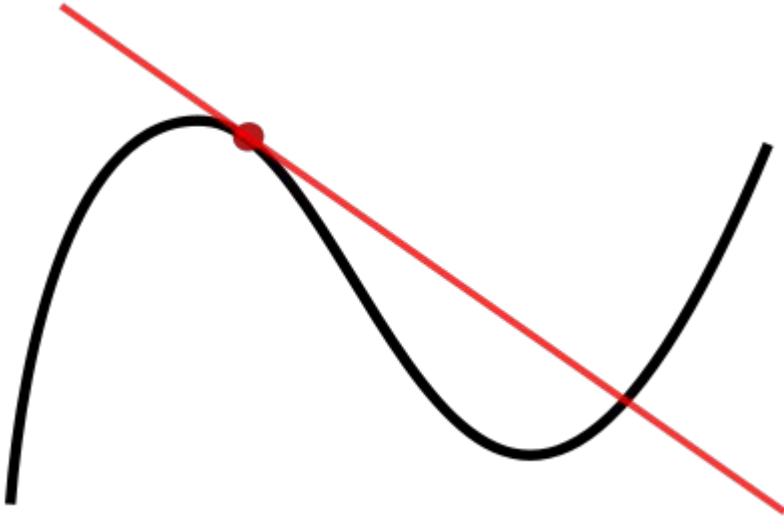
Partial Derivative

Derivatives calculated on vectors (gradient)

Line has slope $\dfrac{\partial f}{\partial x}(a, b)$

Graph of f(x,b)

Point (a,b,f(a,b))
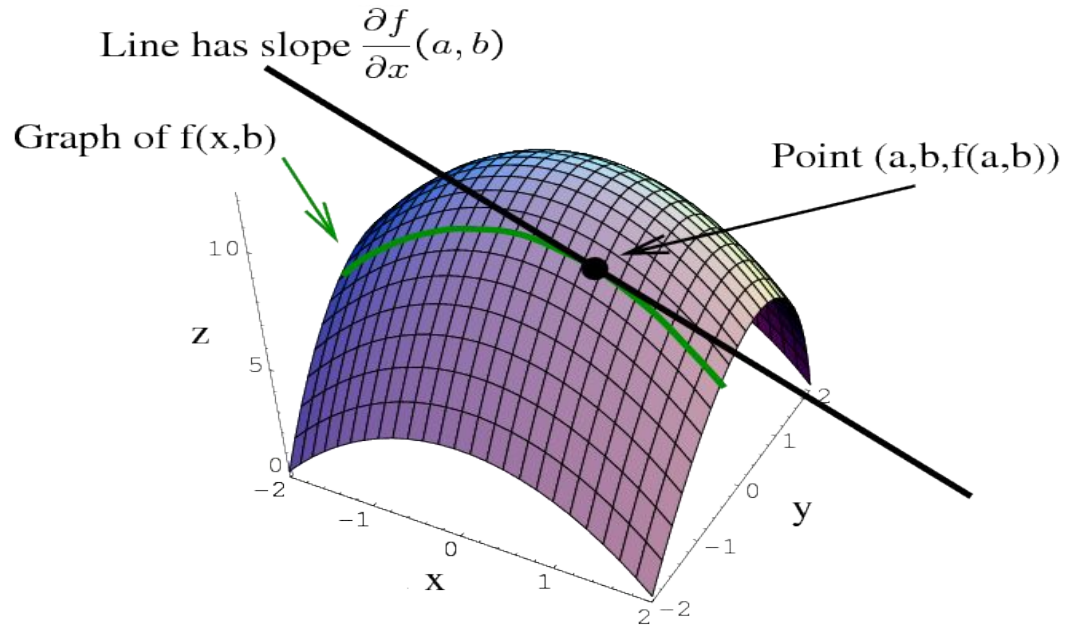
Jacobian Matrix

Used to take derivatives of vectors (arrays), aka gradient

$$\frac{\partial(x,y)}{\partial(u,v)} = \begin{vmatrix} \dfrac{\partial x}{\partial u} & \dfrac{\partial x}{\partial v} \\ \dfrac{\partial y}{\partial u} & \dfrac{\partial y}{\partial v} \end{vmatrix} = \frac{\partial x}{\partial u}\frac{\partial y}{\partial v} - \frac{\partial x}{\partial v}\frac{\partial y}{\partial u}$$

Hessian Matrix

.... is the 2nd derivative of the vector, the gradient of the gradient

used to find maxes (negative),

mins (positive), saddle points

(mixed signs)

$$\text{Gradient} = \begin{bmatrix} \dfrac{\delta H}{\delta z_1} \\ \dfrac{\delta H}{\delta z_2} \\ \dfrac{\delta H}{\delta z_3} \end{bmatrix}, \text{Hessian} = \begin{bmatrix} \dfrac{\partial^2 H}{\partial z_1^2} & \dfrac{\partial^2 H}{\partial z_1 \partial z_2} & \dfrac{\partial^2 H}{\partial z_1 \partial z_3} \\ \dfrac{\partial^2 H}{\partial z_2 \partial z_1} & \dfrac{\partial^2 H}{\partial z_2^2} & \dfrac{\partial^2 H}{\partial z_2 \partial z_3} \\ \dfrac{\partial^2 H}{\partial z_3 \partial z_1} & \dfrac{\partial^2 H}{\partial z_3 \partial z_2} & \dfrac{\partial^2 H}{\partial z_3^2} \end{bmatrix}$$

Constrained Optimization

Used to force RNNs, Swarms, SVM to a global minimum solution of a convex problem

Maximize f(x,y) subject to g(x,y)

…. possibly h(x,y), j(x,y)……

Lagrangian

Example: Solve for both Revenue and Cost

R(x,y) = lambda * C(x,y)   # 3 unknowns (lambda, x, y),   need 3 equations

Take gradient and set equal to zero

L(x, y, lambda) = R(x,y) - lambda * C(x,y)

L(x, y, lambda) = 0

Lagrange Multiplier and Constrained Optimization

https://datastoriesweb.wordpress.com/2017/06/25/lagrange-multiplier-and-constrained-optimization/

Karush-Kuhn-Tucker Approach

- Combining Lagrange conditions for equality and inequality constraints yields *KKT conditions* for general problem:

$$\min_{\mathbf{x}} \quad f(\mathbf{x})$$
$$\text{s.t.} \quad \mathbf{g}(\mathbf{x}) \leq \mathbf{0}$$
$$\mathbf{h}(\mathbf{x}) = \mathbf{0}$$

Lagrangian:

$$L = f(\mathbf{x}) + \boldsymbol{\mu}^T \mathbf{g}(\mathbf{x}) + \boldsymbol{\lambda}^T \mathbf{h}(\mathbf{x})$$

$\Rightarrow$

$$\frac{\partial L}{\partial \mathbf{x}} = \frac{\partial f}{\partial \mathbf{x}} + \sum \mu_i \frac{\partial g_i}{\partial \mathbf{x}} + \sum \lambda_i \frac{\partial h_i}{\partial \mathbf{x}} = \mathbf{0} \quad \text{(optimality)}$$

and

$$\mathbf{g} \leq \mathbf{0}, \quad \mathbf{h} = \mathbf{0} \quad \text{(feasibility)}$$

$$\boldsymbol{\lambda} \neq \mathbf{0}, \quad \boldsymbol{\mu} \geq \mathbf{0}, \quad \mu_i \, g_i = 0 \quad \text{(complementarity)}$$

http://slideplayer.com/slide/4991139/

Machine Learning Basics

Chapter 5 Deep Learning

Resources

Bay Area Deep Learning School 2016

https://www.youtube.com/watch?v=eyovmAtoUx0 and https://www.youtube.com/watch?v=9dXiAecyJrY

Deep Learning Summer School Montreal

https://www.youtube.com/playlist?list=PL5bqIc6XopCbb-FvnHmD1neVlQKwGzQyR

http://videolectures.net/deeplearning2017_montreal/

Neural Networks for Machine Learning (Hinton Lectures 78 videos)

https://www.youtube.com/playlist?list=PLoRl3Ht4JOcdU872GhiYWf6jwrk_SNhz9

Must Know Tricks in Deep Neural Networks

http://lamda.nju.edu.cn/weixs/project/CNNTricks/CNNTricks.html

The Task

This is what you want your neural network to solve.

Think carefully, you can get very different answers depending on how you frame your task.

-- First self driving cars were trained with human drivers and learned to drive on their own.

-- Test car is successfully driven down the road, then down a hill into a lake

? Car learned to follow the curb

Classification

Classification: image recognition, medical symptoms, marketing, voters ....

Features are turned into vectors and the distance between vectors is used to determine group membership. (city block, bird's eye)

Features representing values are scaled between 0,1 or -1,1

Features representing classes are converted to one hot vectors

Classification is unsupervised learning and might not always classify in a useful format. Re-evaluate your features if it's not returning useful classes

Classification with missing inputs

Joint Training of Deep Boltzmann Machines for Classification, Goodfellow

https://arxiv.org/pdf/1301.3568.pdf

Deep Boltzmann Machines, Salakhutdinov and Hinton

http://proceedings.mlr.press/v5/salakhutdinov09a/salakhutdinov09a.pdf

Deep Learning 4j has a nice tutorial on Boltzmann machines

https://deeplearning4j.org/restrictedboltzmannmachine

(aka Deep Belief Networks)
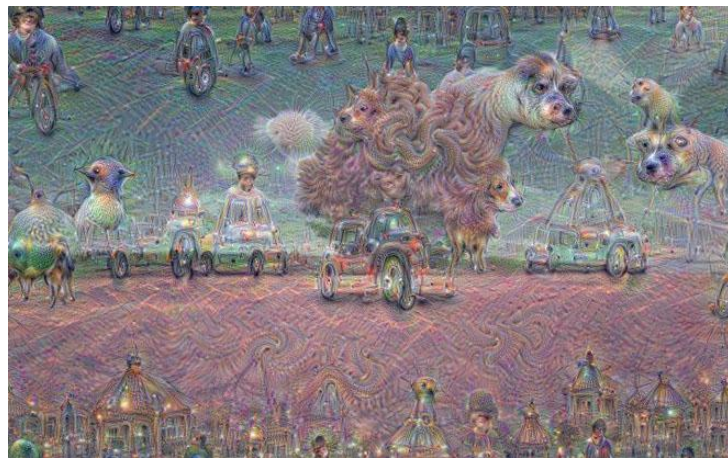
Linear Least Squares Regression

Solve for unknown f(x)

$J(w) = \|y - Xw\|^2$

Use derivative to find direction to change weights

dJ/dw = 2 * XTXw -2 * XTy

w = w - learning_rate * 2 * (XTXw - XTy)

Transcription

Only saw one recent reference to transcription, a deep CNN that crops, segments and outputs street numbers

Multi-digit Number Recognition from Street View

https://static.googleusercontent.com/media/research.google.com/en//pubs/archive/42241.pdf

Goodfellow talk on the paper

https://www.youtube.com/watch?v=vGPI_JvLoN0

Dataset

http://ufldl.stanford.edu/housenumbers/

Machine Translation

Used for natural language translation

Typically done with a recurrent network (RNN, LSTM, GRU, ...)

One language is the input, the second language is the output. The network is trained back and forth rather than just forward feeding data and back feeding the errors

https://nlp.stanford.edu/

https://github.com/MicrosoftTranslator/DocumentTranslator

Anomaly Detection

Make a prediction, points outside the prediction area are anomalies

RAD - Outlier Detection on Big Data

https://medium.com/netflix-techblog/rad-outlier-detection-on-big-data-d6b0494371cc

Robust PCA

http://statweb.stanford.edu/~candes/papers/RobustPCA.pdf

Bayesian Anomaly detection

http://www.dtic.mil/cgi-bin/GetTRDoc?AD=ADA610860

The Performance Measure

~ 80% of data is used for training

~ 10% of data to validate

~ 10% of data is held out to be used as a test after training is done

While accuracy less than (pick a number)%:

      run training data in batches
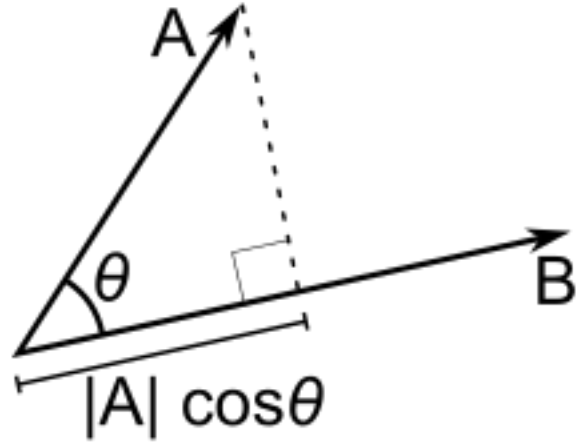
      run validation every 100th epoch or so:

Then test on hold out data after training is complete

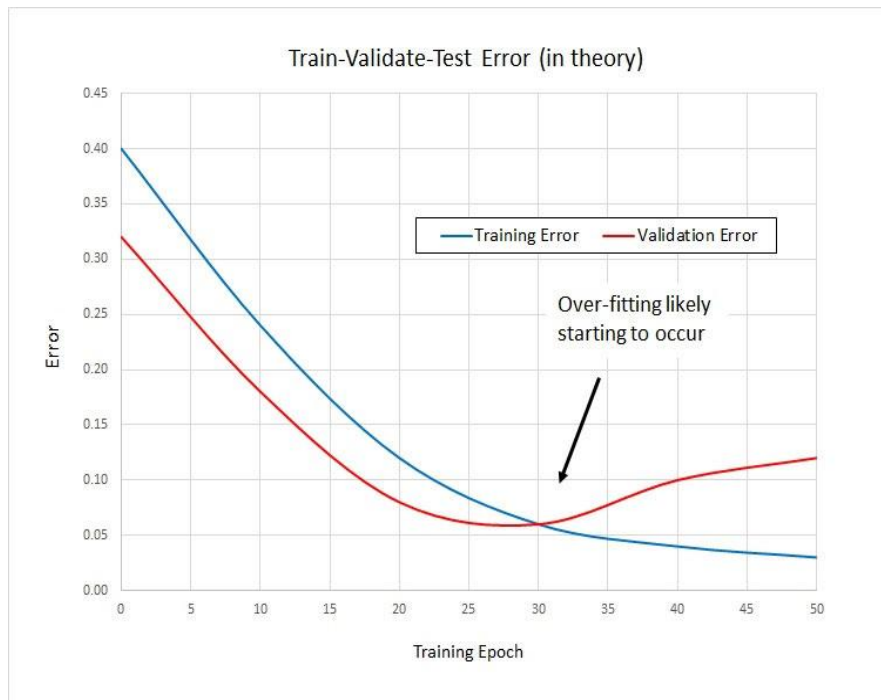Training, testing, validation data must be statistically similar

Linear Regression

y_hat = dot(w.T, x)

MSE = 1/m * Sum(y_predicted - y_test)^2
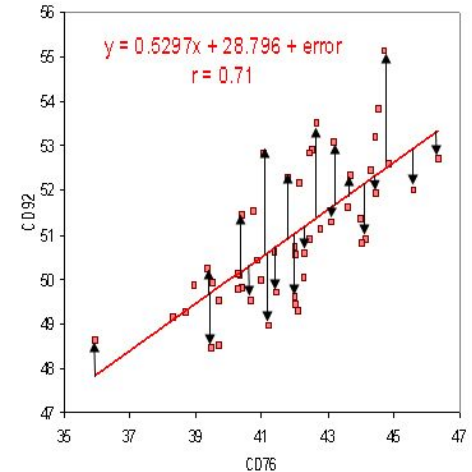
Capacity, Overfitting, Underfitting

Train-Validate-Test Error (in theory)

Training Error — Validation Error

Over-fitting likely starting to occur

Confusion Matrix - very useful for debugging

No Free Lunch Theorem

Says that if you average over all the machine learning algorithms each will have the same performance on the same data set.

The only way one strategy can outperform another is by adjusting the network to the problem.

- convolutional networks for vision
- recursive networks for time series

http://www.asimovinstitute.org/neural-network-zoo/

Hyperparameters

Variables used to control the algorithm

- max, min, depth of branches on a decision tree
- number of layers, size of layers in network

Usually a grid search (nested loops over several options) is used with cross validation

http://scikit-learn.org/stable/modules/grid_search.html

A recent development is that randomly picking values for a grid search is more effective than stepping through all values.

http://www.jmlr.org/papers/volume13/bergstra12a/bergstra12a.pdf

Bias vs Variance

High bias happens during under fitting, high variance during overfitting

Standard error is the square root of variance
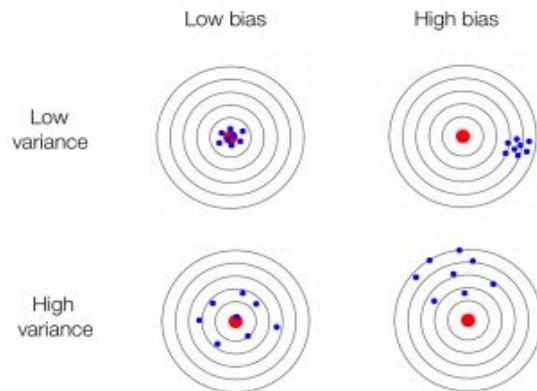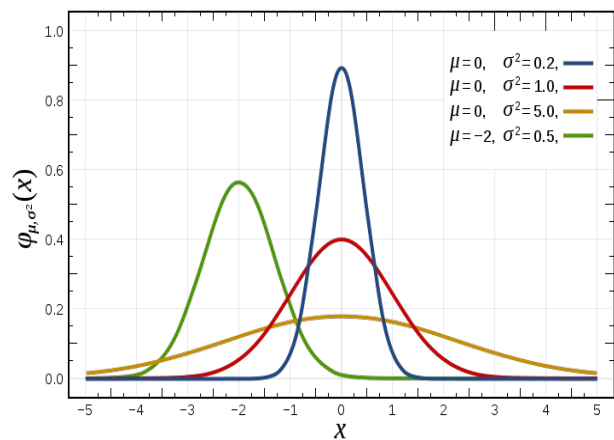
Bias is hitting the 12 on the dart board every time, variance is hitting a wide area on the board

Entropy

Entropy = how many bits are needed to encode information

(* information theory uses Log2, not e, not 10)

$$Entropy = -\sum_{i} P_i \, Log_2 \, P_i$$

A Mathematical Theory of Communication, Shannon

http://affect-reason-utility.com/1301/4/shannon1948.pdf

KL Divergence

Kullback-Leibler Divergence tells how much information is lost between the neural network prediction and the actual values

$$D_{KL}(p||q) = \sum_{i=1}^{N} p(x_i) \cdot (\log p(x_i) - \log q(x_i))$$

The principle of maximum likelihood says that given the training data, we should use as our model the distribution f(w) that gives the greatest possible probability to the training data
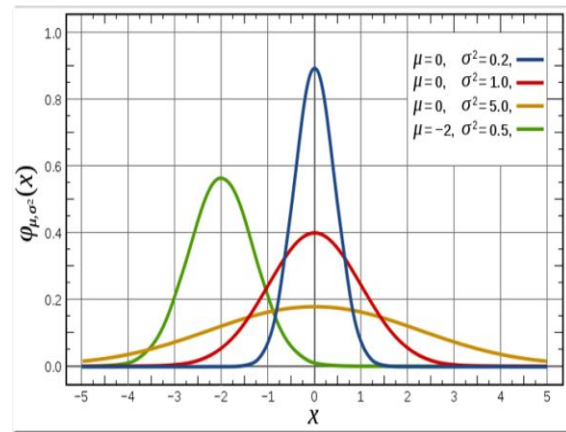
https://www.countbayesie.com/blog/2017/5/9/kullback-leibler-divergence-explained

http://cseweb.ucsd.edu/~elkan/250B/logreg.pdf

Maximum Log Likelihood

$$\theta_{MLE} = \underset{\theta}{\arg\max} \log P(X|\theta)$$

$$= \underset{\theta}{\arg\max} \log \prod_i P(x_i|\theta)$$

$$= \underset{\theta}{\arg\max} \sum_i \log P(x_i|\theta)$$

Conditional Log Likelihood

Weights are trained to maximize the conditional data likelihood

$$W* = \underset{W}{\text{argmax}} \prod_{d \in D} P(Y^d \mid X_1^d ... X_n^d)$$

which is the same as maximizing the conditional log likelihood

$$W* = \underset{W}{\text{argmax}} \sum_{d \in D} \ln P(Y^d \mid X_1^d ... X_n^d)$$

http://cseweb.ucsd.edu/~elkan/250B/logreg.pdf

Mean Squared Error

If the data examples are i.i.d. ...

independent and identically distributed, then the Conditional Log Likelihood can be reduced to the Mean Squared Error

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^{n} (y_i - \tilde{y}_i)^2$$

Independent Identically Distributed

Garbage in, garbage out

Independent

   - features cannot be correlated

Identically Distributed

- training, validation, test data
- outputs (if 90% of the output is true, the network will always guess true and you'll have 90% accuracy - check the confusion matrix)

https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.corr.html

http://scikit-learn.org/stable/modules/generated/sklearn.metrics.confusion_matrix.html

Linear Regression as Maximum Likelihood

Dot product as a measure of similarity between vectors

MSE ~ Conditional log likelihood

As n_samples -> infinity convergence increases

- True distribution of data must be in model family
- True distribution must correspond to one set of weights
-

Let's take the Con out of Econometrics

http://www.econ.ucla.edu/workingpapers/wp239.pdf

Bayesian Statistics

The Posterior	The Evidence	The Prior
	The probability of getting this evidence if this hypothesis were true	The probability of H being true, before gathering evidence

$$P(H|E) = \frac{P(H|E)\,P(H)}{P(E)}$$

The probability that the hypothesis (H) is true given the evidence (E)

The marginal probability of the evidence (Prob of E over all possibilities)

Intuitive understanding of Bayes

https://arbital.com/p/bayes_frequency_diagram/?l=55z&pathId=28771

http://scikit-learn.org/stable/modules/naive_bayes.html

Bayesian Linear Regression

- Prior is uniform or Gaussian distribution with high entropy
- Observation reduces entropy and drives weights to single values

Using a known prior probability

Prior distribution shifts weights toward simpler smoother curves
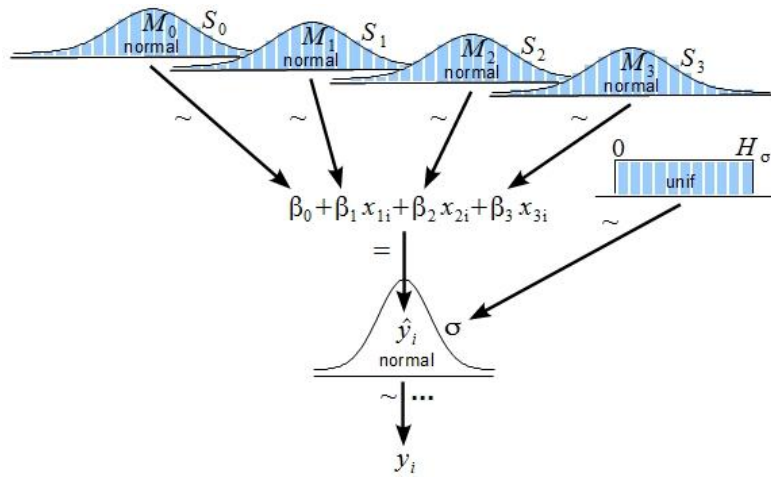
Maximum a Posteriori (MAP) Estimation

$$\theta_{MAP} = \arg\max_\theta P(X|\theta)P(\theta)$$

$$= \arg\max_\theta \log P(X|\theta)P(\theta)$$

$$= \arg\max_\theta \log \prod_i P(x_i|\theta)P(\theta)$$

$$= \arg\max_\theta \sum_i \log P(x_i|\theta)P(\theta)$$

Maximum Likelihood and Maximum Posteriori Estimation

$$\theta_{MAP} = \arg\max_{\theta} \sum_i \log P(x_i|\theta)P(\theta)$$

$$= \arg\max_{\theta} \sum_i \log P(x_i|\theta)\, const$$

$$= \arg\max_{\theta} \sum_i \log P(x_i|\theta)$$

$$= \theta_{MLE}$$

https://wiseodd.github.io/techblog/2017/01/01/mle-vs-map/

Probabilistic Supervised Learning

Normal (Gaussian Distribution)

Probability of y given x

The activation functions are used to create non-linearity

- sigmoid
- tanh
- relu
- softmax (used in last layer to scale values so they sum up to 1)

Bayesian vs Frequentist

Bayesian

Frequentist

Prior beliefs

Repeatable random sample

Fixed parameters

Fixed data

mu = mx + b

y = mx + b

y = N(mu, std)

N - normal distribution

std - standard deviation

mu - mean

https://stats.stackexchange.com/questions/252577/bayes-regression-how-is-it-done-in-comparison-to-standard-regression
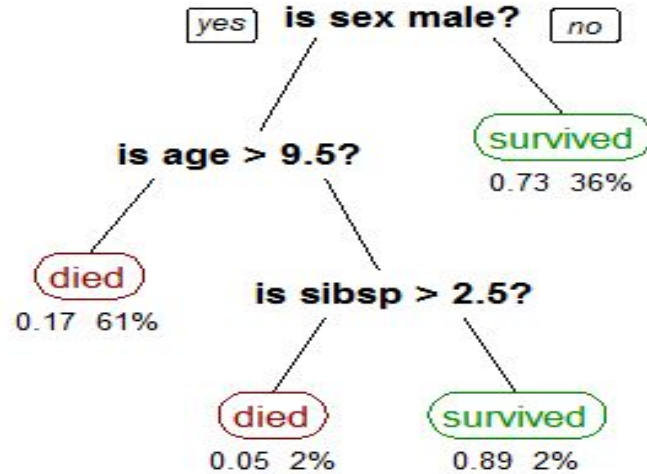
Support Vector Machines
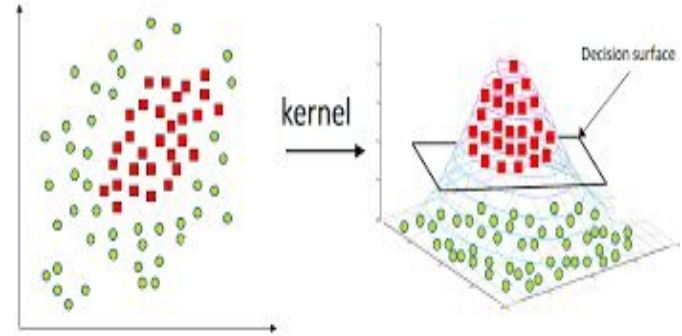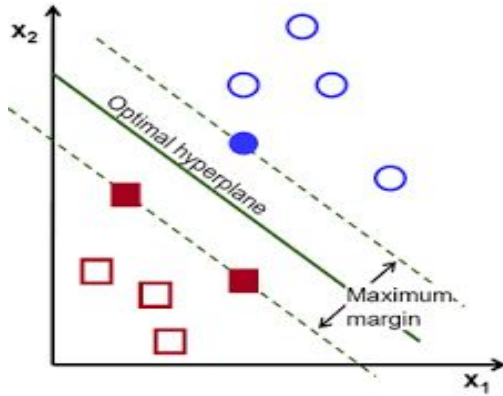
O(n^2), can only separate 2 classes per SVM

http://scikit-learn.org/stable/modules/svm.html

https://docs.opencv.org/2.4.13.4/doc/tutorials/ml/introduction_to_svm/introduction_to_svm.html

Decision Tree

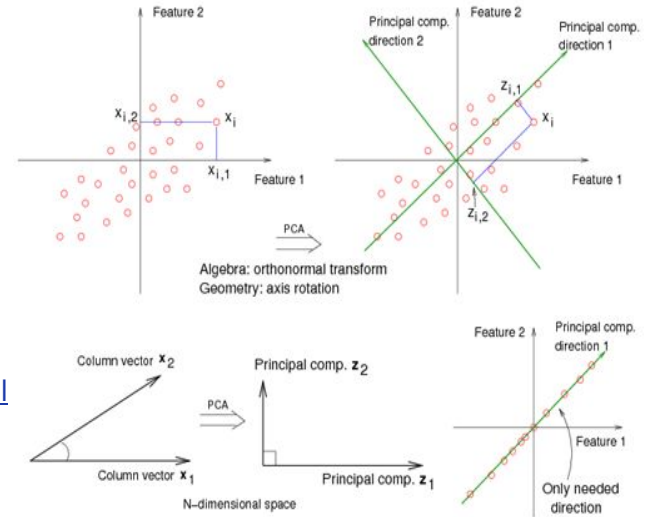Using decision trees to see what a neural network is doing

Principal Components Analysis

Used to compress data to reduce features by combining them in a way that maximizes variance

Use this when data is highly correlated to create independent feature representations

https://onlinecourses.science.psu.edu/stat857/node/35

http://scikit-learn.org/stable/modules/generated/sklearn.decomposition.PCA.html

Time Series Predictions

Used to predict future values

- Must be stationary (rotate down to x axis)
- Generally use a log scale

Facebook's Prophet is a great open source tool for doing forecasts with large amounts of data
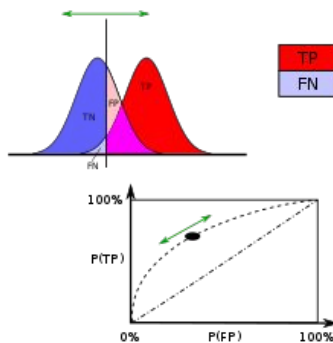
https://facebook.github.io/prophet/

K-Means Clustering

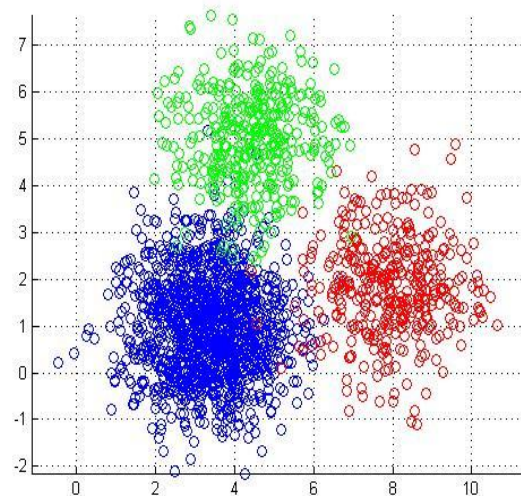pick n points at random (centroids), must give algorithm the number of clusters

put each point in the group of the closest centroid using either Manhattan or Euclidean distance

ROC Receiver operating characteristic ~ 1

http://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html

http://stanford.edu/~cpiech/cs221/handouts/kmeans.html
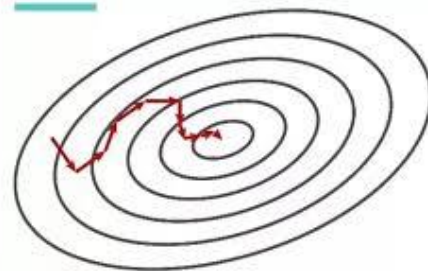
Stochastic Gradient Descent

Repeat until convergence

{

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^{m} (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

}

Stochastic Gradient Descent

Building a Machine Learning Algorithm

Cleaned data

Multiply it by weights, iterative training for non-linear data

Minimize difference between cost function and actual data

(* no free lunch theorem)

New things being tried (or re-tried)

NEAT, NeuroEvolution of Augmenting Topologies

Genetic Algorithms

Reinforcement Learning

Curse of Dimensionality

This is the same issue you've seen in linear algebra.

If you have 2 equations and 3 unknowns there are an infinite number of solutions.

If you have a lot of features and few training samples there will be many paths through a network, they may validate and test okay, but will often randomly fail in real world use.

Regularization

J(w) = Mean Standard Error + regularization

Training adjusts the weights to minimize the error, regularization forces the weights to stay small which forces the network to generalize

Used to prevent overfitting: To be useful a network must be able to generalize to new data.

- L1 regularization
- L2 regularization
- Early stopping
- Drop out
- Batch normalization

L1 Regularization

Minimize the sum of the absolute differences

Drives some weights to zero which has the effect of pruning features which acts as a feature selector. Almost never used in real problems

L1 = c1 * sum(abs(W))

L2 Regularization

Least squares: Minimize the square of the difference

Limits weight size without driving any to zero

Rotational invariance is the main reason it is used instead of L1

L2 = c2 * sum(W)^2

Early Stopping

Plot results during training

Stop training when training and validation errors converge

Dropout

Averages predictions over all weights

Randomly exclude a percent of weights from the neural network, change and randomly exclude a different group on next batch during training.

Dropout: A Simple Way to Prevent NN from Overfitting

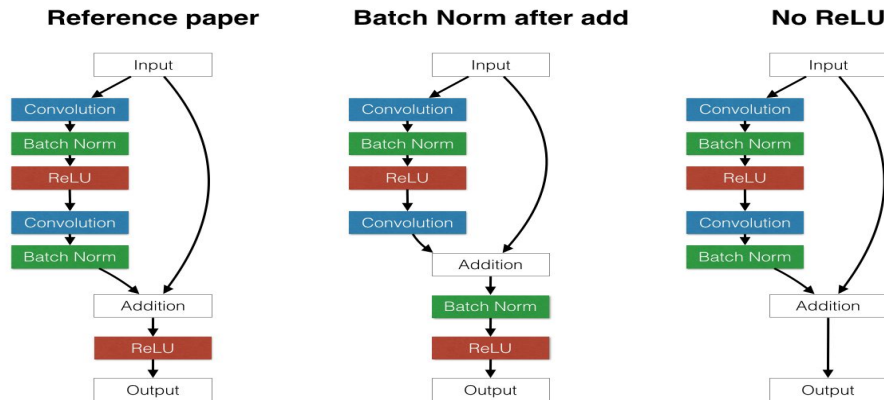https://www.cs.toronto.edu/~hinton/absps/JMLRdropout.pdf

Batch Normalization

Is often used in place of DropOut, can use higher learning rates, networks train faster.

For each batch of training examples recenter the features.

In practice, the original input is re-fed to later layers

| Reference paper | Batch Norm after add | No ReLU |

https://arxiv.org/pdf/1502.03167.pdf

Manifold Learning

Find a useful way to represent the data in less dimensions

Example: represent road as 2D in a 3D world

It's like PCA but attempts to retain more information

t-SNE Used often for visualization

Doesn't always work, information can be lost or distorted

http://scikit-learn.org/stable/modules/manifold.html